

COMP591 Research Paper

Analysing Branch Prediction Algorithms using
OpenPAT



The University of Waikato

Neil Bradley

Proforma

Author: **Neil Bradley**
Project Title: **Analysing Branch Prediction Algorithms
using OpenPAT**
Supervisor: **Dr Simon Spacey**
Word Count: **8100**
Version: 17 February 2015
Contact: ncb9 AT students.waikato.ac.nz

NON DISCLOSURE AGREEMENT

This work contains information deemed commercially sensitive to the University of Waikato, and covered under an agreement of non-disclosure. Please do not distribute or discuss without prior permission. Reading implies acceptance of these terms.

Contents

| | |
|--|-----------|
| List of Figures | v |
| List of Tables | v |
| 1 Introduction | 1 |
| 2 Background | 5 |
| 2.1 Prediction Mechanisms | 5 |
| 2.1.1 Static Prediction | 5 |
| 2.1.2 Profiled Information | 6 |
| 2.1.3 Saturating Counters | 7 |
| 2.1.4 Pattern History Tables | 7 |
| 2.1.5 History Registers | 8 |
| 2.1.6 Branch Target Buffers | 9 |
| 2.1.7 Two-level Predictors | 9 |
| 2.1.8 Combined Predictors | 10 |
| 2.1.9 Predictor Sizes | 11 |
| 2.2 Benchmarking Methods | 12 |
| 2.2.1 Benchmark Suites | 12 |
| 2.2.2 Profiling Tools | 12 |
| 2.3 Performance Metrics | 13 |
| 2.3.1 Success Rate | 13 |
| 2.3.2 Failure Rate | 13 |
| 2.3.3 Contiguous Execution | 14 |
| 2.3.4 Average vs Total vs Best | 14 |
| 3 Implementation | 15 |
| 3.1 Predictor Implementation | 15 |
| 3.1.1 Configuration Testing | 16 |
| 3.1.2 Optimizing | 16 |

| | | |
|----------|--|-----------|
| 3.1.3 | Biased Predictors | 17 |
| 3.1.4 | Combinations | 17 |
| 3.2 | OpenPAT Configuration | 19 |
| 3.2.1 | Target Address | 19 |
| 3.2.2 | Bias Bits | 19 |
| 3.3 | MiBench Configuration | 20 |
| 3.4 | Branch Densities | 20 |
| 3.5 | Branch Offset Distribution | 21 |
| 4 | Results | 23 |
| 4.1 | Best Predictors per Benchmark | 24 |
| 4.2 | Overall Predictors | 27 |
| 4.3 | Top 5 Predictors for Sizes | 29 |
| 4.4 | Branch Density | 31 |
| 4.5 | Jump Distance | 31 |
| 5 | Evaluation | 33 |
| 5.1 | Branch Distance Considerations | 34 |
| 5.2 | Processor Performance from Predictor Accuracy and Branch Density | 35 |
| 5.3 | Restrictions of Study | 36 |
| 6 | Conclusion | 39 |
| | Bibliography | 41 |
| | Appendices | 45 |
| A | Sample Predictor Code Listing (gshare) | A1 |
| B | Sample Predictor Totals (gshare) | B1 |
| C | Totals Tables | C1 |

List of Figures

| | | |
|-----|---|----|
| 1.1 | Pipelined CPU Stages and Execution | 2 |
| 2.1 | Alternative Saturating Counter State Machines | 7 |
| 2.2 | A Simple Two-Level ‘ <i>local</i> ’ Predictor | 9 |
| 2.3 | McFarling’s local/gshare combined predictor | 10 |

List of Tables

| | | |
|-----|------------------------------------|----|
| 3.1 | Branch Predictors Tested | 16 |
| 3.2 | Benchmarks Used | 20 |
| 4.1 | Branches Per Instruction | 31 |

Acknowledgements

I am grateful to my supervisor, Dr. Simon Spacey, whose help and guidance was invaluable in the completion of this project. To Brad Cowie and the Waikato University Computer Science department, for the provision of dedicated workspace. To Assoc. Prof. David Bainbridge, who provided tips and hints along the way. And finally, to my patient (ie, long-suffering) wife Sophie - life is just simply better with you.

Chapter 1

Introduction

Humans are obsessed with a need for speed. This can be seen in the popularity of all sorts of races, but nowhere is it more evident than in the field of computer science. Since the first digital computer ENIAC in 1946 [2], scientists, mathematicians and engineers have produced ever-increasing speeds of calculation, always pushing the bounds of what is physically possible for the production techniques of the day. From an initial speed of 5kHz, over the last seven decades processor speed has developed to over 4GHz in production processors, with CPUs over 10GHz in development [1,2].

A number of technological advances made this progress possible: First, computers went from valve powered to transistor to integrated circuits. As integrated circuits matured, die sizes became ever smaller, allowing more processing power and faster transmission times between components. At the same time, processors were partitioned into task-units, or *pipelined*.

Processing an instruction is a multi-step process - instructions need to be fetched, decoded, executed, and then the results written back to either memory or the next instruction. In a pipelined processor, each of these tasks, and/or their sub-tasks, is assigned to a single unit, running for one clock cycle and passing its state on to the next unit for the next processor cycle. Overall, this results in each instruction taking longer to process (as it now takes multiple cycles), but allows multiple instructions to be processed concurrently, each using a different task-unit. Because each task-unit is now physically much smaller than the whole processor was before, this partitioning also allows processors to operate at much higher clock speeds than before. But it also brings with it a significant problem, in the form of data dependencies.

A data dependency is where one instruction relies on the result of a previous instruction or instructions in order to make its calculation. But if the instruction depended on is still in the pipeline, unresolved, the dependent instruction is

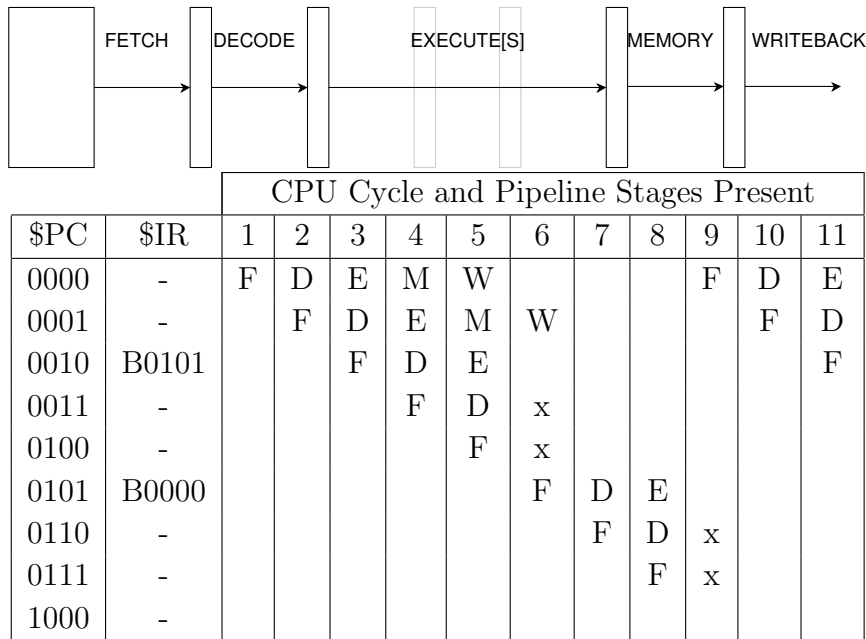


Figure 1.1: Pipelined CPU Stages and Execution

unable to proceed. For most instructions, say, calculations, this is readily solved by both forwarding unresolved instructions and looking ahead deeper into the pipeline to get results early. However, there is one class of instruction for which these options are not enough - the conditional jump, or branch, instruction.

The problem posed by branch instructions is that they alter the processors fetch cycle. Where the processor's program counter or instruction register is merely incremented for a normal instruction, a branch instruction alters the very value of the register, depending on either the condition being tested or the value of a register. But in a pipelined processor, this value may not yet be available, meaning that the processor is unable to process the branch instruction in the current fetch cycle. This leaves the processor with two different options: first, it could *stall* the processor, that is wait doing nothing until the condition is resolved (this could also be achieved by the insertion of `nop` instructions). Or, the processor could attempt a guess at the branch outcome - a technique known as *branch prediction*.

Each of these options is not without it's detriment, however. Inserting `nop`'s or otherwise stalling the pipeline reduces the processors efficiency, by reducing the amount of instructions it can process in parallel. And as branches can make up to an estimated thirty percent of processed instructions, even a shallow pipeline will suffer a significant performance penalty [22]. The other alternative,

branch prediction, suffers no penalty when the prediction is correct. However, if incorrect, it now suffers the same performance degradation of stalling, with the added overhead of having to flush the pipeline - *squashing* the invalid instructions and ensuring they do not corrupt state [22]. Figure 1.2 illustrates such a processor, and a fictional program experiencing these stalls. Careful examination shows this fictional program squashes two out of every five instructions, and therefore only 60% of the work the CPU is doing is actually useful.

A stalling pipeline is much easier to implement than one which utilises a branch predictor. But as we have already seen, a stalling pipeline guarantees reduced performance, whereas a branch predictor promises the potential for increased performance, depending on the accuracy of the predictor utilised. In our quest for ever-increasing speed, therefore, it is desirable to have pipelined processors with highly accurate branch predictors.

Over the last four decades, many predictor designs have been proposed to tackle this need-for-speed. This research project will investigate some of their performances, and methods by which their performance may be improved. Section 2 investigates the history of branch predictor designs, including componentry, algorithms, and performance testing. Section 3 explains the method used for the testing of a selection of branch predictors, some modifications investigated, benchmark and performance analysis configuration, and the rationale behind choices made. Section 4 presents the results of the investigation, with Section 5 discussing relevant findings and limitations of the study, before conclusions and recommendations are presented in Section 6.

Chapter 2

Background

2.1 Prediction Mechanisms

Branch predictors can only make predictions based on the information available to them at the time of execution. There are six main sources of information that can be used: rules, branch target address, program counter, histories, profile information, and the `OP_CODE` itself. What makes branch predictors more successful or not is the algorithms each uses to combine these sources of information into a prediction. In pipelined processors, the mis-prediction penalty is usually directly related to pipeline depth. Hwu *et al* note that, if the penalty for an incorrect prediction is high enough, even a small increase in the accuracy of a branch predictor has a large effect on the performance of the processor [18].

2.1.1 Static Prediction

Early methods of branch prediction made simple guesses about generic branch behaviours based on perceived programmer habits. Depending on how a programmer for a processor was recommended to order branches, the options for the predictor were `ALWAYS_TAKEN`, `NEVER_TAKEN`, `BACKWARDS_TAKEN`, or `FORWARDS_TAKEN`. The latter two options used the branch target address to determine the action to be taken, whilst the former used no information at all. `ALWAYS_TAKEN` is beneficial in applications that involve a large amount of loops, for example, numerical processing, as loops are generally taken far more than not [22]. `NEVER_TAKEN` is beneficial in applications involving large amounts of `switch`-type instructions, as most of the branches will not be taken. Good examples of this kind of behaviour are regular expression evaluators and text sorts. `BACKWARDS_TAKEN` and `FORWARDS_TAKEN` married these two behaviours together, as both behaviours are usually present in the same application. Predictions made

in this way have maximum benefit if the program is written/compiled for them. Programmer habits such as writing most likely outcome in the corresponding branch entry stem from these behaviours. Emer and Clark report a 67% accuracy for an `ALWAYS_TAKEN` predictor on the VAX 11/780 [14].

2.1.2 Profiled Information

But what if the compiler or programmer doesn't program in the way the predictor expects? Or there are branches that, due to certain conditions, go the opposite way? In such cases, the predictor is now more likely to be wrong than right, and the processor suffers performance degradation.

One method implemented to counter this involves the use of extra instruction bits to indicate the branches behaviour [24,25]. In 1986, McFarling and Hennessey suggested that predictor performance could be further increased if the compiler knew which branches are usually not taken. They suggest using a 'squashing direction' bit to convey this information to the branch predictor [22]. Hwu *et al* referred in 1989 to a 'prediction bit' in the branch instruction format, which is set by profiling. They call this method the *Forward Semantic*, and suggest a two-stage compilation with the first stage using nodes inserted at block entries to gather data on branch behaviour [18]. Wall found static prediction based on a profile to perform almost as well as a simple dynamic predictor (bimodal) [32].

These *hint* bits are generated using a two-stage compilation [10]. The first stage compiles the program and runs it with representative data. From this run, the program is then tagged, with the bias bits either stored in the branch instruction, or the programs data section, depending on the architecture. The advantage of this scheme is that if necessary, the program can be re-profiled and the executable re-tagged without a recompilation. The architecture can even be designed in such a way as to re-profile each execution. This is particularly useful for programs where branch behaviour is heavily data dependent.

Predictors that use either profiled information or the `OP_CODE` are typically classified as static predictors [10,33]. Smith first proposed such a predictor in 1981 to save the calculation time of testing branch direction. Several variants have been proposed since, with several authors proposing "rules" that should be followed - the most notable example of which is the paper "Branch Prediction for Free" by Ball and Larus [4].




Figure 2.1: Alternative Saturating Counter State Machines

2.1.3 Saturating Counters

Of course, it is not always possible to know when a branch will be taken. Many branches' behaviour is data-dependent, for which a static prediction may not be suitable. Therefore a second class of predictor, *dynamic* predictors which make predictions based on run-time data were proposed. At the heart of most such predictors lies a simple predictor - the saturating counter [29]. Patented by J. Smith in 1979, saturating counters are a means by which a processor may be 'self-profiling' [29]. The most common version uses two bits to store information about branch direction. If the last branch was taken, the counter is incremented, without overflow. If the last branch was not taken, the counter is decremented, without underflow. A prediction can be made from the higher order bit of the counter - if it is a 1, the branch is predicted **TAKEN**, otherwise it is predicted **NOT TAKEN**. Two variations of this are shown in Figure 2.2. The middle two modes, 1 and 2, are known as **WEAKLY_NOT_TAKEN** and **WEAKLY_TAKEN** respectively. Nair concluded that out of 5248 possible variations of the 2-bit saturating counter, the simple up-down variant performs the best overall [25]. 1-bit and 3-bit counters are also sometimes used/recommended, but the 2-bit is the most prevalent [29].

2.1.4 Pattern History Tables

By itself, a single saturating counter doesn't make much of a predictor. Therefore, it is common practice to have a set of saturating counters, known as a *Pattern History Table* (PHT), which is indexed via an algorithm. PHTs typically make up the bulk of the silicon of a branch predictor, with larger tables increasing predictor accuracy. Predictors often have multiple PHTs, and the subtlety of each predictor is in the algorithm used to select which saturating counter to use to make the prediction with. Early predictors, such as the *bimodal* predictor, used the branch instruction address to index into the table [22, 29]. Each byte

2.1.5 History Registers

For example, say there is a program matching the letter ‘a’ from the string “abababababababababababababab”. At first glance, it looks like the predictor will be right 50% of the time, but this is not necessarily the case. The accuracy of the saturating counter to predict `TAKEN` or `NOT_TAKEN` actually depends on the initial state of the counter. It just so happens that, if the counter is at `WEAKLY_NOT_TAKEN` when the first ‘a’ is encountered, it will always be wrong. This is because the counter will adjust in the opposite direction whenever it is wrong, but by starting at the `WEAKLY` opposite, it always changes away from what would be the next correct prediction. Thus, in the scenario given, a saturating counter will be correct 50% of the time *at most*, and may possibly even be 100% wrong.

Consider the example again, this time using a history register to index the saturating counter. As the program progresses, the local history looks something like “10101010101010101”. If even only two bits of history register are used, whenever the program is about to process an ‘a’, the register contains the value ‘10’, but when it’s about to process a ‘b’, the register contains ‘01’. It can quickly be calculated that the saturating counter indexed by ‘10’ will quickly become **TAKEN**, and the counter indexed by ‘01’ becomes **NOT-TAKEN**. Barring the

few initial mis-predictions based on relative counter state, the predictors are now always correct, and the previous 50% maximum accuracy no longer holds.

History registers also have the advantage that they are simple to implement, use very little hardware, and the next prediction can be known well in advance [5, 26]. Nair suggests using the `BRANCH_TARGET` to generate a history instead of `TAKEN / NOT_TAKEN` [24]. Programs with fewer active branch sites, which are more difficult to predict due to higher data dependencies, benefit more from longer histories than from increased PHT capacity [6].

2.1.6 Branch Target Buffers

Waiting for a predictor to make a prediction takes time. In order to alleviate this, and thus speed up the processor, many processors use a cache known as a *Branch Target Buffer* (BTB) to speed up selection of the target [5]. A branch target buffer stores the most likely target, hashed against the branch address, for quick evaluation upon fetch. Using this BTB, it is possible to reduce branch fetches to a single cycle [29]. Indirect branches cannot be so easily predicted as conditional jumps, so they are predicted using BTBs instead [32]. In order to facilitate faster predictions, Yeh and Patt suggest also storing the next prediction for a given branch address in the BTB when the branch predictor is updated with the correct prediction status [33, 34]. Calder and Grunwald argue against BTBs, arguing that profiled predictors consistently outperform any form of dynamic predictor [5].

2.1.7 Two-level Predictors

In the previous example, a local history was used. In order for such a local history to exist, however, the predictor must have the capacity to hold multiple branch histories, necessitating the use of multiple tables. In 1992, Yeh and Patt proposed using two levels of tables in order to generate a prediction: the first was a *Branch History Table* (BHT) which held local histories, and was indexed by the program counter; the second was a Pattern History Table, and was indexed by the Branch History Table (see Figure 2.3). Yeh and Patt showed this particular design (later called a *local* predictor [21]) to saturate around 97% accuracy, though they used averages of the SPEC benchmarks to claim this [33]. Most predictors since then have been multiple levels.




Figure 2.2: A Simple Two-Level ‘local’ Predictor

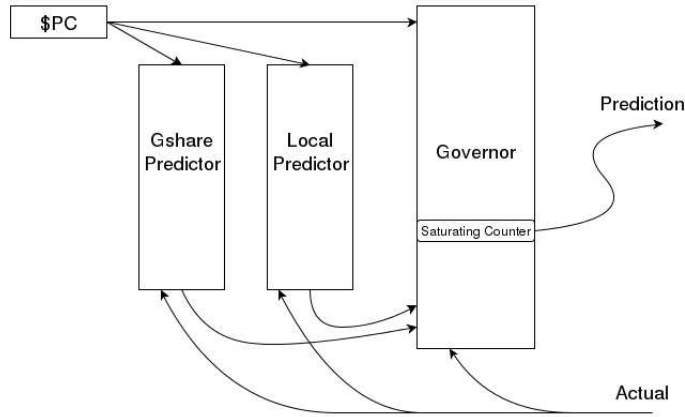


Figure 2.3: McFarling's local/gshare combined predictor

2.1.8 Combined Predictors

In 1993, Scott McFarling of DEC proposed another way to improve predictor accuracy - *combining* predictors. This involves taking two (or more) smaller predictors and combining them with a selector mechanism in order to choose which predictor to use for the current branch. McFarling's design centered around using a table of saturating counters for each branch, thus keeping track of which predictor was most accurate for the branch selected (see Figure 2.4). He proposed and tested two designs: the *local/gshare* and *local/bimode*, and claimed a saturated accuracy of 98.1% [21].

Other methods have also been proposed for combining predictors, such as the *Cascaded* predictor and static methods. In his 2000 PhD paper ‘Improving Branch Prediction by Understanding Branch Behaviour’, Marius Evers argued

against static methods for predictor combination, stating that “the best predictor for a branch changes between data sets. This indicate[s] that it is probably best to use a dynamic selection mechanism to select which component predictor to use for each branch” [15]. Whilst some authors have argued against combined predictors, most predictors since McFarling’s work have been combined predictors [13].

Evers *et al* proposed a combined predictor that used a selection of accuracy counters to choose which predictor to use for a branch. In their predictor, all sub-predictors generate predictions, and the accuracy counters utilise a priority encoder to select which predictor makes final selection. They claimed 97.13% accuracy for their benchmark suite at 64KB [16].

Grunwald *et al* propose using a profiled hint bit to select the best predictor to use, instead of a dynamic selector [17]. Young *et al* suggest using profiling to assign dynamic predictors to difficult-to-predict branches (‘per-stream’), and leaving easy-to-predict branches to a static predictor [36].

2.1.9 Predictor Sizes

One of the most commonly used ways to improve predictor performance has been to increase the size of the tables underlying the algorithm. Larger PHTs mean more saturating counters available to be assigned to pattern predictors, and larger BHTs mean more varied histories available to be assigned. The reason increasing each of these capabilities is important is to reduce *collisions* in predictor selection. However, larger predictors don’t always result in more accuracy. Yeh and Patt noted that increasing history length rather than PHT size resulted in better performance. Their view may be coloured however, by their inclusion of floating point benchmarks with very few branch sites [35]. Chang *et al* discussed this in relation to McFarling’s *gshare* predictor, where they showed that, despite the large silicon area assigned to the predictor, most of it remained unused for their benchmark suite, due to the high number of collisions involved [7].

Larger predictors also consume more power and execute slower. Most early predictor designs were optimistic in the amount of future space allocable for their predictors. Chang *et al* proposed 32K predictors in 1994, for example, and Evers tested predictors up to 384KB in size [8, 15]. Processor speeds were still increasing rapidly, advancements in manufacturing techniques were allowing more silicon space (in accordance with Moore’s Law), and power was relatively cheap. Approximately 5% of a modern CPU is dedicated to branch prediction [11]. Since the mid-2000’s, however, there has been a trend to smaller, more power-efficient predictors (and processors). Therefore, an important trade-off to consider in predictor design is the size of the predictors.

2.2 Benchmarking Methods

Over the years, the methods of testing branch predictor designs has varied. When evaluating an author's claimed accuracy, it is necessary to take into account the benchmark they used to test this. Because of the variance in methods used, it is hard to compare published results directly - rather most authors include previously published predictors in their implementations to facilitate comparison. There is, however, an alarming trend to ignore more recent predictors; most authors compare their work with McFarling's *gshare* predictor (which XORs the `BRANCH_ADDRESS` and `GLOBAL_HISTORY` to index its PHT) only, despite it being superseded in his own work. Early implementers also used to restrict the running time of their benchmarks to a certain number of instructions - understandable given the cost of computing power at the time [21].

2.2.1 Benchmark Suites

When branch prediction was first introduced by researchers, they typically used small Unix programs, such as `wc` and `gzip` to test predictor performance. Later, the SPEC benchmark suite became popular [4, 5, 8, 21]. Early researchers often used the full suite, however, since 1994 it has been the trend to use a reduced set, as certain benchmarks were shown to skew the results due to high predictability. The primary culprits here were floating point benchmarks, and certain others that had highly used loops which were not data dependant [24]. Five of Yeh and Patts original ten benchmarks, for example, were floating point, and conversely, exhibited less increase in accuracy as predictor sizes were increased - though were very accurate initially [33, 35]. A more recent benchmark often used is the IBS Ultrix traces, which includes OS calls, library calls and context switches, and substantially more branches than SPEC [10]. Sechrest *et al* note that more recent benchmark suites tend to have a higher percentage of difficult branches [28].

2.2.2 Profiling Tools

Various methods have also been used to get the information required for testing. Most authors run their tests on profiled information, either gathered themselves or obtained from a third party source. Dynamic testing has also been done - that is, the calculations made during execution. One such program is `pixie`, used by Smith *et al*: '`pixie` is a program that allows you to trace, profile, or generate dynamic statistics for any program that runs on a MIPS processor. `pixie` is fast and convenient since it annotates executable object code with additional

instructions that collect the dynamic information during run time. Your program is run directly on the machine, no emulation is done nor are additional system calls made.’ [30] Calder *et al* used ATOM, a two-staged tool that profiled then executed their benchmarks [5]. Sprangle *et al* used SoSS, a virtual system with a profiling API, which ran at 1/30th of a normal CPU [31]. And some used QPT on DEC machines, which also generates traces for system libraries [4].

Predictors algorithms themselves have typically been implemented as either stand-alone C programs, or plug-ins for these tools.

2.3 Performance Metrics

As well as using different benchmarks to test their predictors, predictor performances are usually reported against predictor size. However, there is also large variance in the way authors report their findings. The main methods of reporting are **success%**, **failure%**, and *contiguous instruction count*. Some authors have used more esoteric methods (such as the ratio between profiled accuracy and actual accuracy [27]), but these are usually individual papers. What is most noticeable is that authors report methods which show their predictor is more accurate than whatever they have chosen as baseline. When combined with the variance in benchmarks used, this makes direct comparison problematic. Young *et al* go so far as to say such comparison is pointless [36].

2.3.1 Success Rate

Success rate is the most basic performance indicator - how many branches were predicted correctly? This is useful for calculating processor efficiency. Pan, So and Rahmeh used prediction accuracy as their chosen metric [26]. However, as predictors get more accurate, the gains are smaller and so this statistic seemingly fell out of favour.

2.3.2 Failure Rate

As predictor designs became more accurate, it became more commonplace for authors to report their performance based on *mis-predictions*. This metric enables people to make claims such as “this predictor has 20% less mis-predictions than *gshare*”, making it look like their predictor makes significantly better improvements than it does. However, this is smoke-and-mirrors, as the reality is these accuracy increases result in very minor speed increases. Combined with the tendency to ignore more recent predictors, this metric makes it difficult to

give credence to the reported results. Yeh and Patt used this method, claiming mis-prediction was the most relevant metric to use as it gives the squash rate [33]. It is interesting to note, however, that in later work they reverted to success rate as the relevant metric [35].

2.3.3 Contiguous Execution

Often there is also the option of reporting *average contiguous execution* - how many instructions are executed without a mis-prediction. There are two problems with this metric, however. First, it is easy to skew with benchmarks that have a large amount of loops. Second, unless the author specifically states branches only, it is easy to skew with benchmarks that have low branch density. As such, this metric does not provide an accurate depiction of predictor performance.

2.3.4 Average vs Total vs Best

There is also variance in how predictor performance is reported in relation to the benchmark suite run. It is common practice to report the *average* performance of the predictor across the benchmark suite, however, this is not a fair representation of how the predictor would actually run in reality. The reason for this is due to the varied complexity and length of the benchmarks used. Small, easily predictable benchmarks have an undue influence on the averages reported. Usually, authors report the individual results for benchmarks, allowing the reader to gauge the accuracy of the predictor for themselves.

Chapter 3

Implementation

In this project, branch predictors were tested using OpenPAT and MiBench. First, each predictor was tested for each size, ranging from 1B through to 1MB, with all possible configurations tested, in order to find the most accurate predictor. Second, the best settings were used to generate per-branch tables for the benchmarks. Third, a profiled(*biased*) predictor was investigated using these tables. And finally, a *Tagged* Branch Predictor(TBP) was investigated using these tables. The research was carried out on LUNACY, a 4GHz Intel i7 with 32GB of RAM and an SSD RAID, running Debian Wheezy.

3.1 Predictor Implementation

A predictor consists of a single C file, with a consistent API. A separate program was used to feed this API with information it would have available at run-time, sourced from trace files generated using a new OpenPAT tool, `branch_trace`. The information made available to the predictors was `BRANCH_ADDRESS`, `BRANCH_TARGET_ADDRESS`, `BRANCH_BIAS` and `BRANCH_TAKEN`. `BRANCH_TAKEN` allows the predictor to update itself after making its prediction. It should be noted that the use of trace files was merely to speed up processing time - OpenPAT can be configured to run the branch predictors directly during profiling. An original tool, `branch_prediction` was used in this way to test the functionality of predictors. A sample predictor (`gshare`) can be found in Appendix A.

Yeh and Patt propose a complex rule for specifying the associated hardware cost of a branch predictor [34]. However, in the same way that the overhead associated with a CPU cache is not specified, I do not include control hardware part of the hardware cost for comparative purposes. Also, like Ball and Larus, amongst others, I do not consider indirect branches [4].

3.1.1 Configuration Testing

Each predictor is initially configured in such a way that it tests all variants possible for the given predictor size. 17 different predictors were tested in this project: 4 static, 10 dynamic and 3 combined; with thousands of different configurations tested. Table 3.1 shows these predictors, and their respective features.

Table 3.1: Branch Predictors Tested

| Predictor | Static | Saturating Counter | PHT | History | Combined |
|---------------------------------|--------|--------------------|-----|---------|----------|
| all [29] | ✓ | | | | |
| backwards [29] | ✓ | | | | |
| forwards [29] | ✓ | | | | |
| bias [18] | ✓ | | | | |
| saturating [29] | | ✓ | | | |
| agree [31] | ✓ | ✓ | ✓ | ✓ | |
| bimode [20] | | ✓ | ✓ | ✓ | |
| markov [10, 23] | | ✓ | ✓ | ✓ | |
| bimodal [21, 29] | | ✓ | ✓ | | |
| global [21] | | ✓ | ✓ | ✓ | |
| local [21, 34] | | ✓ | ✓ | ✓ | |
| gselect [21] | | ✓ | ✓ | ✓ | |
| gshare [21] | | ✓ | ✓ | ✓ | |
| correlated [24] | | ✓ | ✓ | | |
| local/gshare [12, 21] | | ✓ | ✓ | ✓ | ✓ |
| cascaded (all/gshare) [12, 21] | | ✓ | ✓ | ✓ | ✓ |
| cascaded (all/gselect) [12, 21] | | ✓ | ✓ | ✓ | ✓ |

3.1.2 Optimizing

Once all predictor settings had been investigated, the best performing settings for each predictor were filtered out and set into a different per-branch-per-predictor table, which was then used to generate comparative results. These results can be seen in Section 4.1.

3.1.3 Biased Predictors

From the per-branch-per-predictor results, it is possible to calculate how a 1-bit *biased* predictor would behave for the same benchmark. The biased predictor proposed assumes a profiled executable, where one bit from the branch target address is used to indicate whether the branch predictor is accurate or not. If the branch predictor is less than 50% accurate, the bit is set to 1, otherwise 0. The underlying branch predictor's prediction is then muxed with its bias, allowing for a profiled executable to account for predictor inaccuracies. This can be simulated per-branch as follows:

```
if (correct[$PC] < (total[$PC] + 1) / 2) {
    correct[$PC] = total[$PC] - correct[$PC];
}
```

It is essential to note that this process is invisible to the predictors - predictors that update their state based on whether they were correct or not are thus not affected by the bias implementation. Jacobsen *et al* suggested such a predictor in 1996 [19].

3.1.4 Combinations

It is also possible to calculate from the per-branch tables how a *Tagged Branch Predictor* (TBP) would perform. A Tagged Branch Predictor uses a set of predictors which always update every branch, but the predictor used to make the prediction is selected by a *tag bit* in the branch instruction. Once again, this process is invisible to the predictors - they are not effected by the TBP implementation.

The TBP is very similar to Grunwald *et al*'s proposed *Static Hybrid Predictor*, or Chang and Banerjee's *Profile-guided Multi-heuristic Predictor*, except that it allows for more sub-predictors, and every predictor is updated regardless of its selection or not [9, 17]. Evers also considered such a predictor, however he wrote off a profiled version without actually testing it, claiming data variances were too great [15]. This goes against many earlier investigations which showed that in the vast majority of cases, data variance had negligible effect on a profiled predictor, provided the profile data was sufficiently representative of actual use [37]. Chang *et al* also proposed multiple component predictors, however it utilised a single underlying PHT into which all predictors indexed [8]. In the TBP, each predictor has its own inviolable PHT.

A TBP can use as many predictors as bit-space in the instruction allows; for example, a 32-bit instruction with 20-bits of address space could be modified to include two tag bits, with only 18 bits left for address space. It is theorised that this minor trade-off will have next to no effect on software - the largest branch found in preliminary testing used only 14 bits.

In order to find the best combination per size allowance, the following algorithm was used:

```
// Calculating best predictor sets per size
foreach (size) {
    best = {"", "", total=0};
    foreach (subset) {
        foreach (size_combination) {
            total = 0;
            foreach (branch) {
                total += max (subset);
            }
            report (size, subset, size_combination, total);
            if (total > best.total) {
                best = (subset, size_combination, total);
            }
        }
    }
    report (size, best);
}
```

It should be noted this is a simplified version of the algorithm used to calculate combinations. The process is complicated by the various different size combinations possible; static predictors, for example, were calculated as having no space, which allowed more space to be allocated to other predictors in the set. It was also complicated by the fact various predictors already used their own ‘tag-bit’; namely the Agree and other biased predictors. This meant sometimes the subset has less predictors than it’s bit-space would suggest. Preliminary testing also showed that a 3-bit TBP (eight predictors) had insignificant gains over a 2-bit TBP (four predictors) for the selected predictor set, so experimentation was confined to 2-bit TBPs. Across all sizes and optimisation levels, 175,654 different TBPs were tested.

3.2 OpenPAT Configuration

Traces for the benchmarks were gathered using OpenPAT. OpenPAT is an open-source profiling suite developed at Imperial College in conjunction with Cambridge, and is available upon registration from www.openpat.org. OpenPAT uses customised tools in order to generate profiled information - for example, the included `hotspot` tool was used to count the number of IRs in a given trace execution. For this project, results were generated using two custom tools - `branch_predictor` and `branch_trace`. These tools require modifications to the OpenPAT framework to work properly, however.

3.2.1 Target Address

Some of the predictors tested required information not immediately available in OpenPAT's data structures. Therefore, OpenPAT was modified to include in its pre-parsing phase a java program which changed its `.label` field to hold the `TARGET_ADDRESS` for each branch. Due to framework limitations, these are the assembly file line numbers. This could mean that the results may not be identical to the benchmark running on hardware due to changes in the interference pattern. However, as these values are consistent across all benchmark/size tests, their usefulness for comparative analysis is not considered diminished.

Due to the way OpenPAT handles multiple input files, it was also necessary to implement a hashing scheme for both the branch address and branch target address.

3.2.2 Bias Bits

Sprangle claims that using the first encountered direction for the Agree predictor is 85% as accurate as using a profiled version [31]. However, sixteen years earlier, Smith had stated that '[it is possible] to use only the first execution of the branch instruction as a guide; a strategy of this type, although accurate, has been found to be slightly less accurate than other dynamic strategies' [29]. This was found to also be the case for the Agree predictor, so in order to garner best results, my implementation used a profiled bias-bit as its initial indicator.

Yeh and Patt also used the same data set as profiled to benchmark, allowing for perfect testing of the branch predictor[s] [33]. As I am testing optimal settings, it was therefore necessary, for the predictors that use a bias bit to be pre-loaded with a per-branch bias bit. In order to generate these bits, before either the bias or agree predictors were run, a `branch_bias` tool was run first, which determined

the exact bias of each branch. These biases were saved in a separate configuration file, which the respective predictors then loaded during their setup phase.

3.3 MiBench Configuration

MiBench is an benchmark suite designed for testing embedded systems [3]. Its use for profiling was suggested by Dr S. Spacey, author of OpenPAT. Like SPEC, several of the benchmarks contained in the suite were not considered useful for testing branch predictors, as they were far too easy to predict. Benchmarks were run at the highest achievable optimisation level - Table 3.2 records the benchmarks used, their respective optimisation level, and the number of branches encountered in each. In total, the chosen suite contained 485M branches from 4498 sites.

Table 3.2: Benchmarks Used

| Benchmark | Optimisation | Branches |
|--------------|--------------|----------|
| adpcm | O3 | 293M |
| dijkstra | O3 | 48M |
| fft | O3 | 3.6M |
| ispell | O1 | 85M |
| jpeg | O1 | 15M |
| patricia | O3 | 2.9M |
| qsort | O3 | 1.9M |
| rijndael | O3 | 2.4M |
| stringsearch | O3 | 377k |
| susan | O1 | 33M |

3.4 Branch Densities

In order to calculate estimated processor performance, it was necessary to calculate the branch density, or *branches per instruction* (BpI). For this, the branch density per benchmark was found by dividing the total number of instructions run (found using OpenPAT's `hotspot` tool) by the number of branches encountered by `branch_trace`. The results per predictor can be found in Section 4.3.

3.5 Branch Offset Distribution

It was also necessary to calculate the number of bits typically required for a branch instruction, in order to ascertain the feasibility of implementing a TBP for current software. Branch offset was measured on **LUNACY** by disassembling **/bin/***, and parsing the result with **grep** to test branch distance:

```
objdump -d /bin/* | grep 'j' > dump.csv
```

Minimum bits required for jump (including sign bit) were then calculated and collated using LibreOffice Calc. As this is for indicative purposes only, the small file sizes typical of **/bin/** were not considered a major issue regarding jump distances. Results can be found in Section 4.4.

Chapter 4

Results

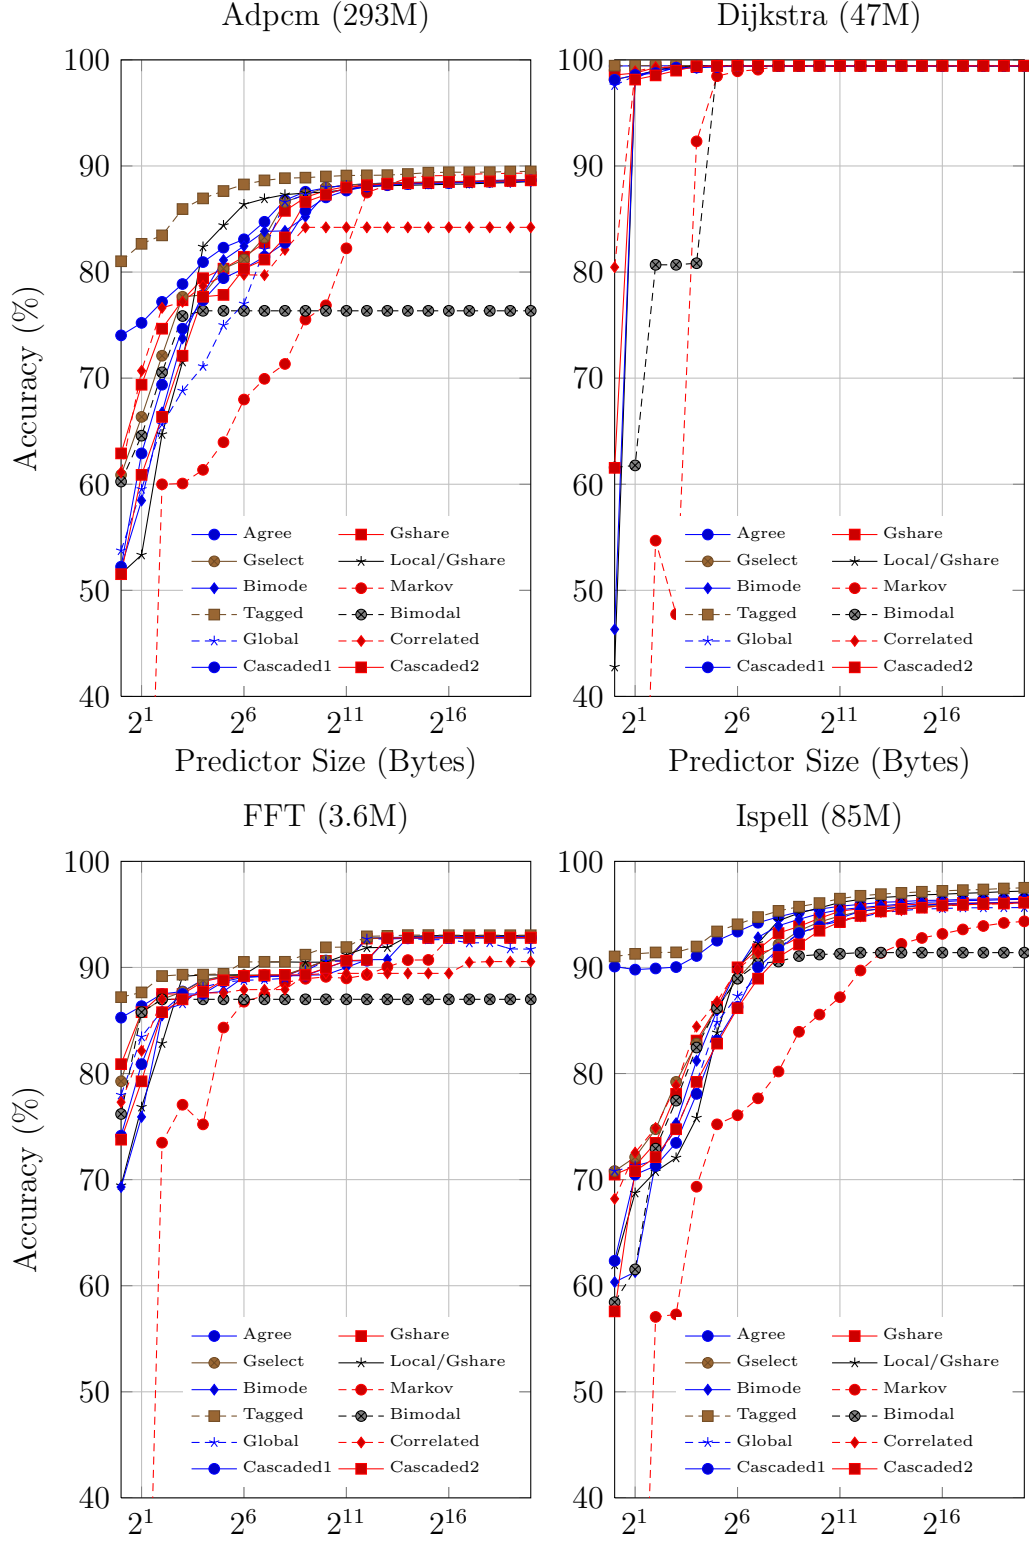
Using the trace files generated by `branch_trace`, two sets of results were garnered for each predictor. The first set, covered in the section Best Predictors per Benchmark, uses the ‘locally’ best settings - that is, the results presented are the best settings for the predictor for the benchmark in question. This means performance between benchmarks is not directly comparable, as the representative predictor may not be the same per benchmark.

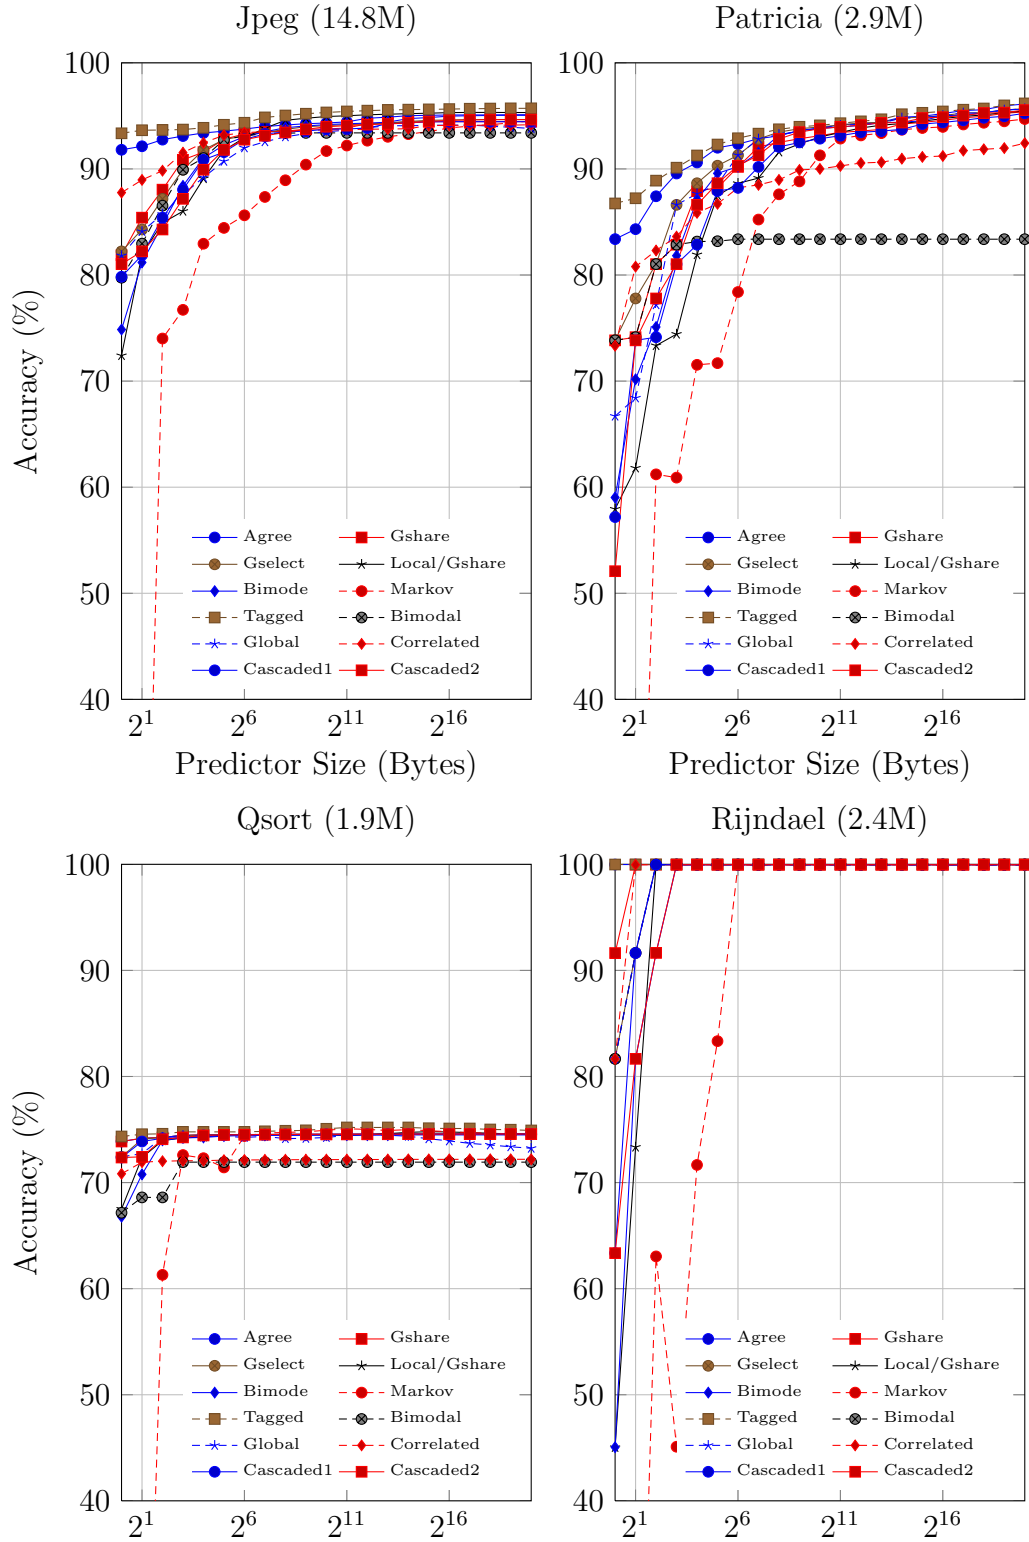
The second set of results contains the overall best performers, both global and local. The ‘global’ best settings are defined as the configurations for each predictor that performed the best overall, according to:

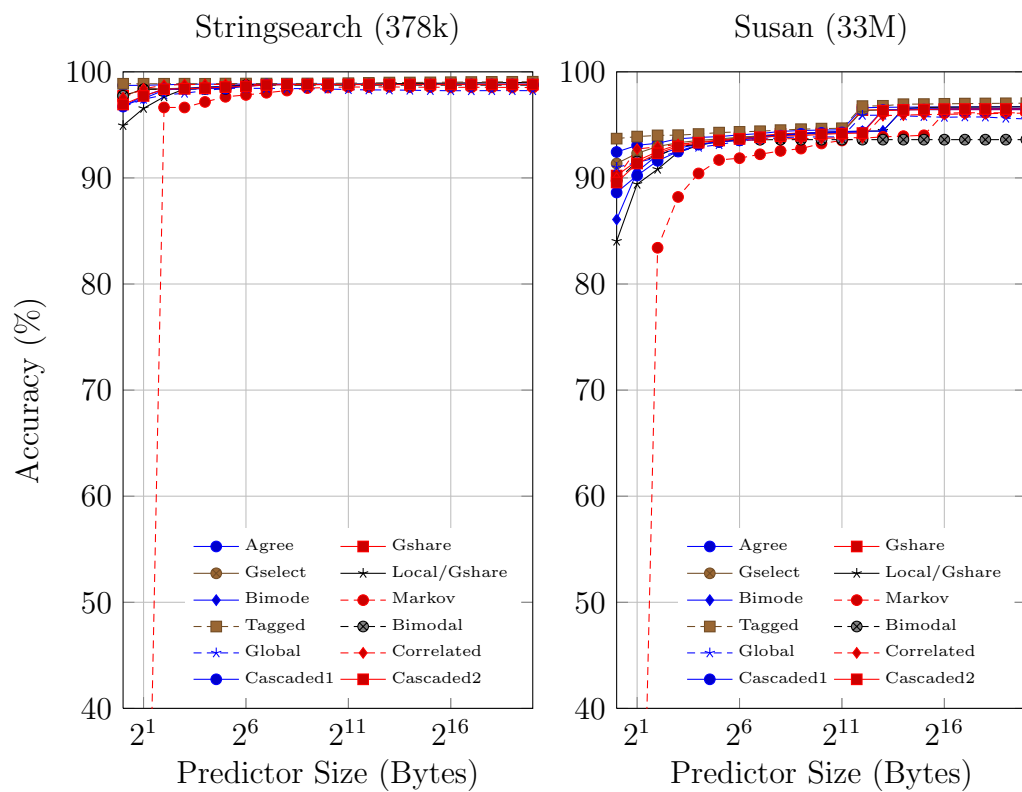
$$\frac{\text{branches_correctly_predicted}}{\text{branches_encountered}}$$

The local best settings uses the best local configurations for each predictor, and takes an unweighted average. The global settings are in effect the same as taking a weighted average, by comparison. Section 4.3 only includes the predictors that were in the top five most accurate for a given size.

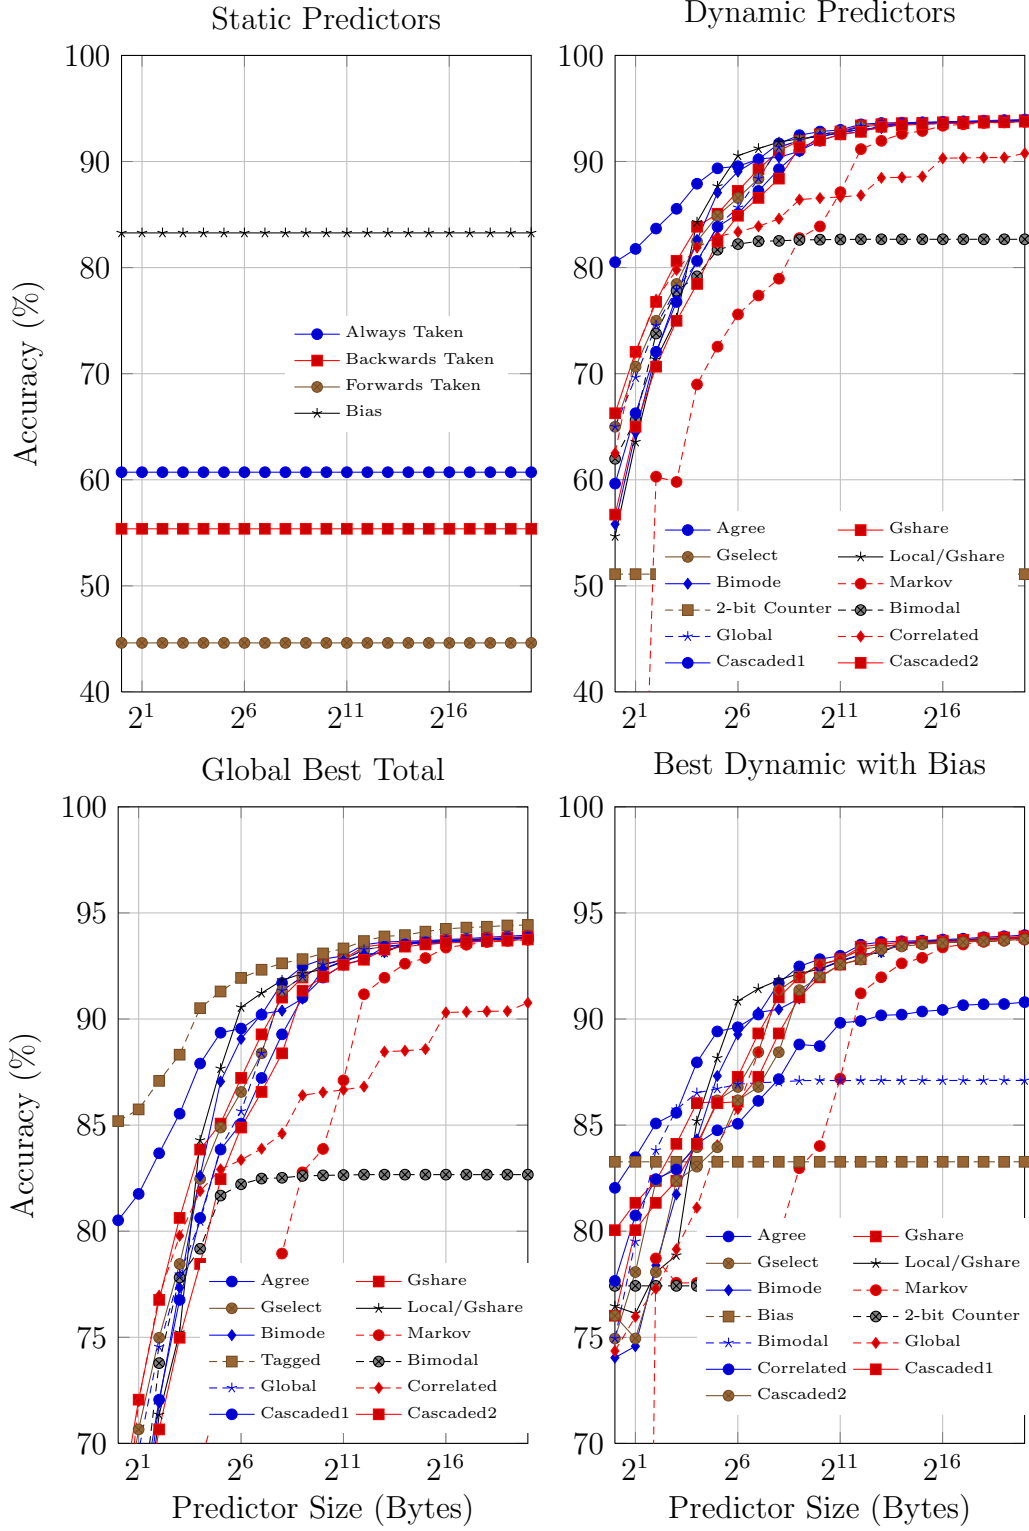
4.1 Best Predictors per Benchmark

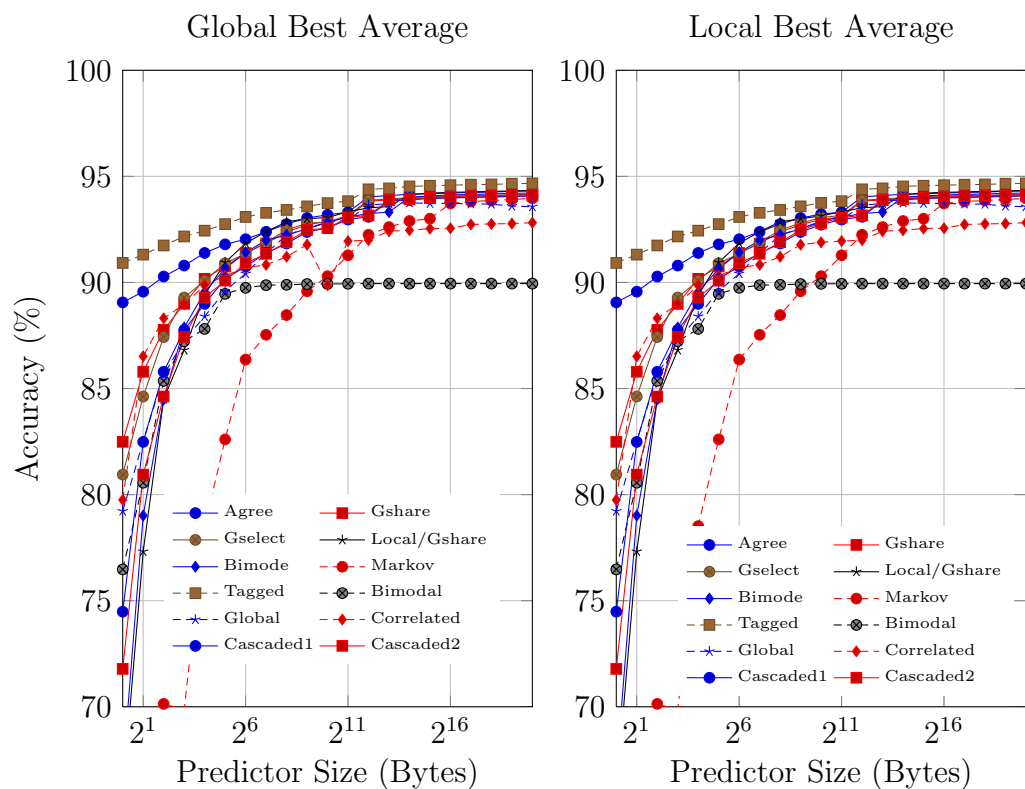







4.2 Overall Predictors






Results tables can be found in Appendix C.

4.3 Top 5 Predictors for Sizes



Top 5 Predictors per Size (Global Best Total)




4.4 Branch Density

Table 4.1: Branches Per Instruction

| Benchmark | Branches | Instructions | IpB | BpI |
|--------------|-----------|--------------|--------|--------|
| adpcm | 293871549 | 1272182080 | 4.32 | 0.2320 |
| dijkstra | 47731814 | 174683579 | 3.66 | 0.2732 |
| FFT | 3596462 | 45261274 | 12.58 | 0.0795 |
| ispell | 85117867 | ² | 6.4 | 0.1562 |
| jpeg | 14806832 | 136502823 | 9.21 | 0.1085 |
| patricia | 2909457 | 20144176 | 6.92 | 0.1444 |
| qsort | 1879057 | 14903057 | 7.93 | 0.1261 |
| rijndael | 2448387 | 407340139 | 166.37 | 0.0060 |
| stringsearch | 376842 | 1572823 | 4.17 | 0.2396 |
| susan | 32680929 | 489816121 | 14.99 | 0.0667 |

4.5 Jump Distance



²The `ispell` benchmark did not finish correctly, causing the `hotspot` tool to not output any results. The figures given here for BpI and IpB are the weighted average of the other benchmarks.

Chapter 5

Evaluation

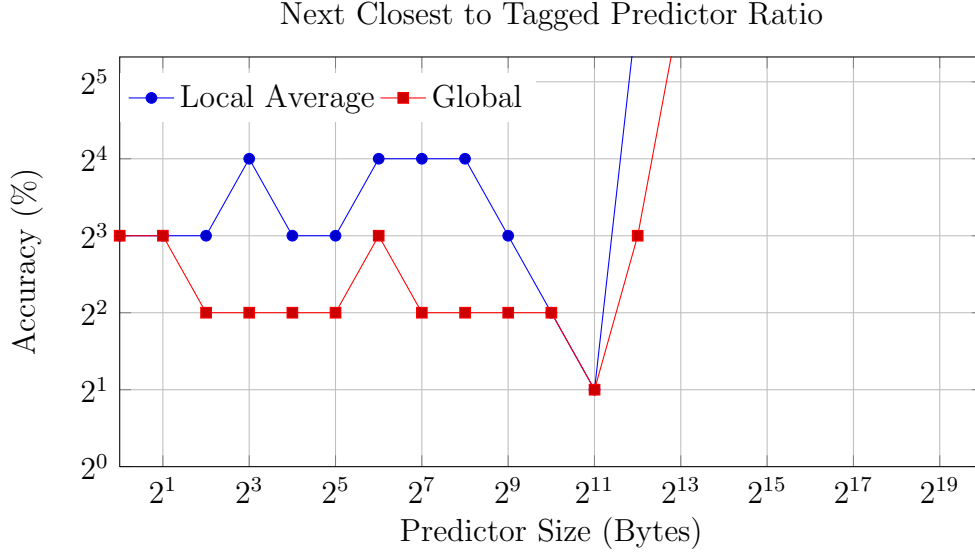
The results obtained in these experiments are not directly comparable to those published by authors. For example, Yeh and Patt reported 97.1% accuracy for their predictor, and McFarling 98.1% for his, and yet here they achieved a best performance of 91.2% and 93.8% at 1MB, respectively. One reason for this is the bias towards difficult-to-predict benchmarks in the MiBench subset, which accounted for almost eighty percent of the benchmark suite.

One very interesting result was the seeming inability of any predictor to accurately predict `qsort`, `fft` or `adpcm`. The reason for this is the same for both benchmarks: `adpcm` has a very small number of branch sites which deal with greatly varying data; likewise `qsort` and `fft`. Other benchmarks exhibiting similar behaviour were `ispell`, `jpeg`, `patricia` and `susan`, though for a different reason: all of these benchmarks had hundreds of active sites against which data is continually tested, making accurate prediction difficult.

Also of interest were the large jumps in accuracy that `susan` and `fft` both experienced around 4KB predictor size. It is unlikely that the predictors accuracy jumping noticeably at 4KB, and the fact that there were 4498 branch sites in the benchmark suite, are unrelated. The reason for this jump is likely to be the reduction in collisions caused by aliasing in the PHTs, and is probably the cause for the jump around 4KB in the total results, though this jump is somewhat smoother in the global total compared to the local average. This is because `fft` only accounted for 0.6% of the branches encountered, and is given unfair weighting in the local average. The jump from `susan` still has a noticeable effect however, since it accounts for approximately 6% of encountered branches.

For predictors that were already very accurate, the biased version produced negligible improvements. However, lower accuracy predictors, notably the static predictors, showed much greater improvements. Having said this, using a bias bit on a static predictor is unnecessary, as it produces the same accuracy as using the

bias bit itself to make the prediction. Nevertheless, at the cost of an instruction bit, biased predictors, as proposed here are not worth the architecture change.



The same can not be said, however, about the Tagged Branch Predictor. It consistently outperformed the other predictors, to such an extent that it outperformed much larger predictors in all except one size bracket (2KB). Figure 5.1 shows the ratio between TBP size and the size required by the next best performer to outperform it. Under 2KB, the TBP uses half, or sometimes a quarter, of the amount of chip space to outperform the next predictor. This has serious ramifications for low-power chips. A 4KB TBP outperforms the very best 16KB predictor tested; and none of the predictors tested at any size, even 1MB, could outperform an 8KB TBP.

The TBP is based on the predictors it was tested against, taking the best of them and combining them into a single predictor. As expected, Agree and local/gshare fared consistently well, and this accuracy was reflected in the TBP. This means that, as even more accurate base predictors are developed, a TBP units performance could be improved by adding it to the subset. For any given predictor size, a TBP will never be out-predicted by a single predictor unit.

5.1 Branch Distance Considerations

The main drawback of the TBP is the architecture change required to implement it - two bits in the branch instruction need to be assigned to be control bits. In a 64-bit system, this is perhaps not a major issue. However, consider a 32-bit

branch instruction, which typically consists of 12 instruction bits, and 20 target bits. If we were to include the TBP control bits in the target section, this reduces the available branch distance to 18 bits, or +/- 128K. Given modern programming paradigms, however, it is theorised that it is unlikely that any branch will exceed this jump distance, and if it were to, this could be got around in an architecture supporting the TBP by using a shorter branch to a jump instruction. This would result in a minor performance impact. This is also not a new proposal: another variant proposed by Calder and Grunwald is the use of indirect branches to facilitate long jump distances [5].

Section 4.5 shows the jump distances of `/bin/*` on the test system `LUNACY`. No branches used more than 18 bits, and in fact, only 6/134217 (0.004%) used more than 15 bits. The impact then, of changing to a TBP on the test system has nothing but positive effects.

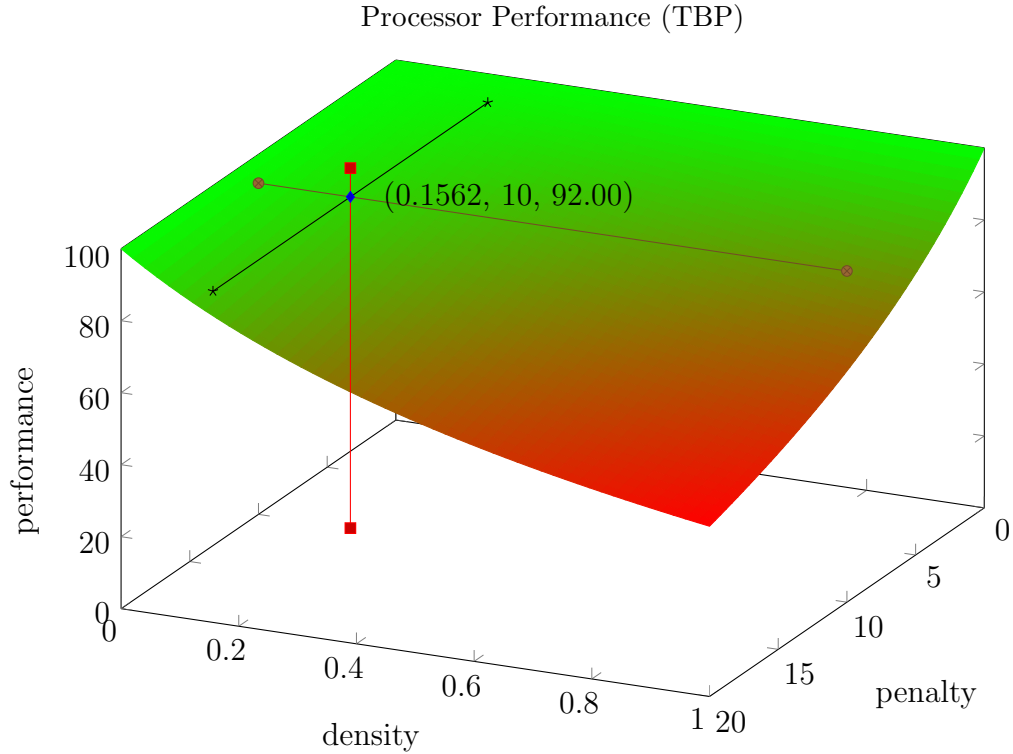
And the TBP provides more advantages than just increased accuracy: For example, any given program can be reconfigured by an external profiler, if required, without a recompile; the TBP will actually execute faster than a combined predictor, due to the single MUX required to implement its selection mechanism; the lower amount of predictor hardware required will engender itself towards low-power chips; the performance increase at larger sizes will have a notable effect on high-end processors; and programs do not have to be profiled on compilation - in this scenario the TBP will merely be a wrapper for the underlying base predictor, which should be the best performer of it's components.

5.2 Processor Performance from Predictor Accuracy and Branch Density

Predictor accuracy, however, doesn't have a direct relationship with CPU performance [37]. Other factors, most notably branch penalty and branch density, also have a significant impact, as seen in Figure 5.2. As pipeline depth increases, due to the associated mis-prediction penalty, processor performance decreases accordingly. The real performance of a processor, as a percentage of its speed, is given by the following formula:

$$perf = \frac{100\%}{1 - density + density * (accuracy + (1 - accuracy) * (1 + penalty))}$$

From this formula, we can calculate that a TBP with an accuracy of 94.5%, on a machine with a 10-cycle penalty (such as the DEC Alpha had [4]), and



running the MiBench sub-suite ($BpI = 0.1562$), will effectively run at 92.00% of maximum stated speed. This predictor has a mis-prediction rate of 5.5%, almost double the 3% Yeh and Patt claimed was unacceptable. Though, their predictor was also tested, and found to have a mis-prediction rate of 8.84% (see Appendix C, “mcf_local”). As has been discussed earlier, this was due to the highly predictable nature of their benchmarks [34]. In any case, their predictor for the same conditions has an effective CPU speed of 87.87% of maximum. In fact, the best predictor tested (a 1MB biased-Agree predictor, which also uses two instruction bits), had an effective CPU performance of 91.37%. This makes a CPU running a TBP 0.6% faster than one running an Agree predictor. In real terms, were the processors 10GHz, that’s an extra 63M instructions processed at no extra cost.

5.3 Restrictions of Study

When looking at the results of this study, it is necessary to keep in mind the restrictions under which the study was done. Many of the criticisms leveled against earlier studies also apply to this one.

First, MiBench hasn’t been updated since 2001, and only a subset of

the benchmark was used. Some elements of the benchmark suite, `rijndael`, `stringsearch`, and `dijkstra` were highly predictable, and account for approximately twenty percent of the branches encountered. The issue this presents is, what level of predictability is acceptable in a benchmark? Whilst it may seem preferable to have a benchmark suite that is as difficult to predict as possible, is this actually indicative of real performance? The answer to this question is beyond the scope of this study, but it bears consideration, as the results of this study are not directly comparable to the results of others for this very reason.

In a similar vein, another restriction is the small set of branches present in the benchmark suite. The MiBench subset used had only 4498 branch sites, which disadvantaged predictors larger than 4KB that did not use any form of history. It also didn't contain any indirect branches, which disadvantaged path-based predictors, such as Nair's Correlated predictor.

Another restriction was that libraries were not instrumented. As noted in Section 2.2.1, several authors claim the results from a single program are not really comparable to real-world application in a multi-tasking system. This is the main motivation behind benchmarks/traces which include system calls and multiple programs. However, whilst this means that the results of this study are not directly comparable to those of other studies, that does not reduce their usefulness for comparative analysis.

The most recent predictor (aside from the TBP) was developed in 1998. There have been other advances in predictor algorithms made since, notably in perceptron-based predictors, and other exotic offerings from the Championship Branch Prediction competition. Nevertheless, regardless of later predictor's performance, they could be easily assimilated into the TBP, and so their relative accuracies are not actually that important.

The final restriction is on the veracity of the results presented. Every possible effort has been made, given the time constraints, to present each tested predictor in its optimal settings. Whilst this was achieved for all the stand-alone predictors, the biased and Tagged predictors were only tested across a subset of their possible configurations. This subset consisted of using the very best settings per component predictor, but it is possible that sub-optimal components may have produced better combined performance. Further investigation may be warranted.

Chapter 6

Conclusion

In this research, optimal settings have been found for a range of predictors. These settings have been used in conjunction with a benchmark suite, consisting of a subset of the MiBench benchmark. The predictors have then been compared for a range of sizes and optimisation levels. A Tagged Branch Predictor has been developed, which uses a small architecture change along with underlying predictor technology to improve branch prediction accuracy. Using a single tag bit to indicate predictor accuracy has also been investigated.

It has been shown that, for the predictors implemented, a biased predictor which uses only a single instruction bit has limited benefit. It has also been shown, however, that the use of more bits for a *tagged* branch predictor results in substantially better predictor performance. If ultimate performance is the aim, the results obtained indicate that a large TBP could be a viable option. Conversely, if the aim is efficiency, a smaller TBP outperforms single predictors over four times larger. Branches larger than that afforded by a TBP, assuming an 18-bit branch offset, are very rare, and could be handled using a short branch to a jump, at very little cost. And there is no cost to using programs either profiled, not-profiled, or re-profiling them either.

This project was also subject to some limitations - particularly in the range of predictors tested. The most recent branch predictor utilised in this study was proposed in 1998. Several more advanced predictors have been proposed since. It would be useful to see how they compare, and affect the TBP implementation proposed.

Bibliography

- [1] Quad core test processor chip doing 10 ghz, August 2013. <https://www.youtube.com/watch?v=MUq50N9w5D8>.
- [2] Eniac - wikipedia, November 2014. <http://en.wikipedia.org/wiki/ENIAC>.
- [3] Mibench version 1.0, March 2014. <http://wwwweb.eecs.umich.edu/mibench/>.
- [4] Thomas Ball and James R Larus. *Branch prediction for free*, volume 28. ACM, 1993.
- [5] Brad Calder and Dirk Grunwald. Fast and accurate instruction fetch and branch prediction. In *Computer Architecture, 1994., Proceedings the 21st Annual International Symposium on*, pages 2–11. IEEE, 1994.
- [6] Brad Calder and Dirk Grunwald. Next cache line and set prediction. In *Computer Architecture, 1995. Proceedings., 22nd Annual International Symposium on*, pages 287–296. IEEE, 1995.
- [7] Po-Yung Chang, Marius Evers, and Yale N Patt. Improving branch prediction accuracy by reducing pattern history table interference. *International journal of parallel programming*, 25(5):339–362, 1997.
- [8] Po-Yung Chang, Eric Hao, Tse-Yu Yeh, and Yale Patt. Branch classification: a new mechanism for improving branch predictor performance. In *Proceedings of the 27th annual international symposium on Microarchitecture*, pages 22–31. ACM, 1994.
- [9] Pohua Chang and Upal Banerjee. Profile-guided multi-heuristic branch prediction. In *1995 International Conference on Parallel Processing*, pages 215–218, 1995.

- [10] J Bradley Chen, Michael D Smith, Cliff Young, and Nicolas Gloy. An analysis of dynamic branch prediction schemes on system workloads. In *Computer Architecture, 1996 23rd Annual International Symposium on*, pages 12–12. IEEE, 1996.
- [11] Karel Driesen and Urs Hölzle. Accurate indirect branch prediction. In *ACM SIGARCH Computer Architecture News*, volume 26, pages 167–178. IEEE Computer Society, 1998.
- [12] Karel Driesen and Urs Hölzle. The cascaded predictor: Economical and adaptive branch target prediction. In *Proceedings of the 31st annual ACM/IEEE international symposium on Microarchitecture*, pages 249–258. IEEE Computer Society Press, 1998.
- [13] AN Eden, J Ringenberg, S Sparrow, and T Mudge. Hybrid myths in branch prediction. In *Proc. 7th International Conference on Information Systems Analysis and Synthesis*, 2001.
- [14] Joel S Emer and Douglas W Clark. A characterization of processor performance in the vax-11/780. In *ACM SIGARCH Computer Architecture News*, volume 12, pages 301–310. ACM, 1984.
- [15] Marius Evers. *Improving branch prediction by understanding branch behavior*. PhD thesis, The University of Michigan, 2000.
- [16] Marius Evers, Po-Yung Chang, and Yale N Patt. Using hybrid branch predictors to improve branch prediction accuracy in the presence of context switches. In *ACM SIGARCH Computer Architecture News*, volume 24, pages 3–11. ACM, 1996.
- [17] Dirk Grunwald, Donald Lindsay, and Benjamin Zorn. Static methods in hybrid branch prediction. In *Parallel Architectures and Compilation Techniques, 1998. Proceedings. 1998 International Conference on*, pages 222–229. IEEE, 1998.
- [18] WW Hwu, Thomas M Conte, and Pohua P Chang. Comparing software and hardware schemes for reducing the cost of branches. In *ACM SIGARCH Computer Architecture News*, volume 17, pages 224–233. ACM, 1989.
- [19] Erik Jacobsen, Eric Rotenberg, and James E Smith. Assigning confidence to conditional branch predictions. In *Proceedings of the 29th annual ACM/IEEE international symposium on Microarchitecture*, pages 142–152. IEEE Computer Society, 1996.

- [20] Chih-Chieh Lee, I-CK Chen, and Trevor N Mudge. The bi-mode branch predictor. In *Microarchitecture, 1997. Proceedings., Thirtieth Annual IEEE/ACM International Symposium on*, pages 4–13. IEEE, 1997.
- [21] Scott McFarling. Combining branch predictors. Technical report, Technical Report TN-36, Digital Western Research Laboratory, 1993.
- [22] Scott McFarling and J Hennesey. *Reducing the cost of branches*, volume 14. IEEE Computer Society Press, 1986.
- [23] Trevor N Mudge, I-Cheng Chen, and John Coffey. *Limits to branch prediction*. University of Michigan, Computer Science and Engineering Division, Department of Electrical Engineering and Computer Science, 1996.
- [24] Ravi Nair. Dynamic path-based branch correlation. In *Proceedings of the 28th annual international symposium on Microarchitecture*, pages 15–23. IEEE Computer Society Press, 1995.
- [25] Ravi Nair. Optimal 2-bit branch predictors. *Computers, IEEE Transactions on*, 44(5):698–702, 1995.
- [26] Shien-Tai Pan, Kimming So, and Joseph T Rahmeh. Improving the accuracy of dynamic branch prediction using branch correlation. In *ACM Sigplan Notices*, volume 27, pages 76–84. ACM, 1992.
- [27] Jason RC Patterson. Accurate static branch prediction by value range propagation. In *ACM SIGPLAN Notices*, volume 30, pages 67–78. ACM, 1995.
- [28] Stuart Sechrest, Chih-Chieh Lee, and Trevor Mudge. The role of adaptivity in two-level adaptive branch prediction. In *Proceedings of the 28th annual international symposium on Microarchitecture*, pages 264–269. IEEE Computer Society Press, 1995.
- [29] James E Smith. A study of branch prediction strategies. In *Proceedings of the 8th annual symposium on Computer Architecture*, pages 135–148. IEEE Computer Society Press, 1981.
- [30] Michael D Smith. *Tracing with pixie*. Computer Systems Laboratory, Stanford University, 1991.
- [31] Eric Sprangle, Robert S Chappell, Mitch Alsup, and Yale N Patt. The agree predictor: A mechanism for reducing negative branch history interference.

- In *ACM SIGARCH Computer Architecture News*, volume 25, pages 284–291. ACM, 1997.
- [32] David W Wall. *Limits of instruction-level parallelism*, volume 19. ACM, 1991.
- [33] Tse-Yu Yeh and Yale N Patt. Two-level adaptive training branch prediction. In *Proceedings of the 24th annual international symposium on Microarchitecture*, pages 51–61. ACM, 1991.
- [34] Tse-Yu Yeh and Yale N Patt. Alternative implementations of two-level adaptive branch prediction. In *ACM SIGARCH Computer Architecture News*, volume 20, pages 124–134. ACM, 1992.
- [35] Tse-Yu Yeh and Yale N Patt. A comparison of dynamic branch predictors that use two levels of branch history. *ACM SIGARCH Computer Architecture News*, 21(2):257–266, 1993.
- [36] Cliff Young, Nicolas Gloy, and Michael D Smith. *A comparative analysis of schemes for correlated branch prediction*, volume 23. ACM, 1995.
- [37] Cliff Young and Michael D Smith. Improving the accuracy of static branch prediction using branch correlation. In *ACM Sigplan Notices*, volume 29, pages 232–241. ACM, 1994.

Appendices

Appendix A

Sample Predictor Code Listing (gshare)

```

/* mcfarling_gshare.c
 * =====
 * By Neil Bradley
 *
 * All rights reserved.
 *
 * This predictor XORs the PC and GHR to index into the PHT.
 *
 * See McFarling, S.; "Combining Branch Predictors", 1993.
 */

#include "predictors.h"
#include <stdio.h>
#include <math.h>

/* MACROS */
#define MCF_GSHARE_PINDEX(a, b)      (((a) * mcf_gshare_variants) + (b))

/* VARIABLES */
int mcf_gshare_size;                // How large the PHTs are
int mcf_gshare_variants;            // How many different variants there are
unsigned int mcf_gshare_GHR;        // The global history register
char *mcf_gshare_PHT;               // [size][variants]
long *mcf_gshare_totals;            // The tallies for each predictor

/* Setup the variables for basic predictor. */
void mcf_gshare_SETUP(int max_size) {
    int i;

    // Calculate sizes
    mcf_gshare_size = max_size;
    i = max_size; mcf_gshare_variants = 1;
    while ((i >>= 1) > 0) {
        mcf_gshare_variants++;
    }

    // Allocate tables
    mcf_gshare_totals =
        (long *)calloc(mcf_gshare_variants, sizeof(long));

```



```

mcf_gshare_PHT =
    (char *)calloc(max_size * mcf_gshare_variants, sizeof(char));

// Initialise totals and PHT
for (i = 0; i < mcf_gshare_variants; i++) {
    mcf_gshare_totals[i] = 0;
}
for (i = 0; i < mcf_gshare_variants * max_size; i++) {
    mcf_gshare_PHT[i] = 3; // Start at fully used.
}
mcf_gshare_GHR = 0xFFFFFFFF;
} // mcf_gshare_SETUP =====

/* The variables passed in provide all the information need to
   calculate the prediction retrospectively. */
int mcf_gshare(unsigned int address, unsigned int target, int taken) {
    int predicted = 0, i, best = 0, prediction = 0, n, m, index;

    // Find best predictor
    for (i = 0; i < mcf_gshare_variants; i++) {
        if (mcf_gshare_totals[i] > mcf_gshare_totals[best]) {
            best = i;
        }
    }
    // Finding best predictor

    // Cycle through predictors
    for (i = 0; i < mcf_gshare_variants; i++) {
        // Calculate parts
        n = address & (mcf_gshare_size - 1);
        m = mcf_gshare_GHR & (~(0xFFFFFFFF << i));
        index = MCF_GSHARE_PINDEX(n ^ m, i);

        // Calculate prediction
        prediction = mcf_gshare_PHT[index] >> 1;
        if (i == best) {
            predicted = prediction;
        }

        // Update PHT

```

```

    if (taken == YES && mcf_gshare_PHT[index] < 3) {
        mcf_gshare_PHT[index]++;
    } else if (taken == NO && mcf_gshare_PHT[index] > 0) {
        mcf_gshare_PHT[index]--;
    }

    // Update totals
    if (prediction == taken) {
        mcf_gshare_totals[i]++;
    }
} // Cycling through predictors

// Update GHR
mcf_gshare_GHR = (mcf_gshare_GHR << 1) | taken;

return predicted;
} // mcf_gshare =====

// Publishes the predictor's current statistics to stdout.
void mcf_gshare_DUMP(void) {
    int i;
    _OP_DISPLAY("gshare (McFarling) predictor active:");
    for (i = mcf_gshare_variants; i > 0; i--) {
        _OP_DISPLAY("\t%d:%d", i - 1, mcf_gshare_totals[i - 1]);
    }
    _OP_DISPLAY("\n");
} // mcf_gshare_DUMP =====

// Cleans up the predictor's memory space.
void mcf_gshare_CLOSE(void) {
    free(mcf_gshare_totals);
    free(mcf_gshare_PHT);
} // mcf_gshare_FINALISE =====

```

Appendix B

Sample Predictor Totals (gshare)

../adpcm_1.trace: Total=146406294 Local={GHR=14b}134471978 LocalInv={GHR=14b_INV}134471978 Global={GHR=14b}134471978 GlobalInv={GHR=14b}134471978
../adpcm_2.trace: Total=133087388 Local={GHR=14b}120686831 LocalInv={GHR=14b_INV}120686831 Global={GHR=14b}120686831 GlobalInv={GHR=14b}120686831
../adpcm_3.trace: Total=7531492 Local={GHR=14b}6475721 LocalInv={GHR=14b_INV}6475721 Global={GHR=14b}6475721 GlobalInv={GHR=14b}6475721
../adpcm_4.trace: Total=6846375 Local={GHR=14b}5769607 LocalInv={GHR=14b_INV}5769607 Global={GHR=14b}5769607 GlobalInv={GHR=14b}5769607
../dijkstra_1.trace: Total=39838870 Local={GHR=4b}39616742 LocalInv={GHR=4b_INV}39616743 Global={GHR=14b}39615129 GlobalInv={GHR=14b}39615130
../dijkstra_2.trace: Total=7892944 Local={GHR=4b}7848955 LocalInv={GHR=4b_INV}7848956 Global={GHR=14b}7848461 GlobalInv={GHR=14b}7848462
../fft_1.trace: Total=1556528 Local={GHR=14b}1420509 LocalInv={GHR=14b_INV}1421853 Global={GHR=14b}1420509 GlobalInv={GHR=14b}1421853
../fft_2.trace: Total=1589296 Local={GHR=14b}1453273 LocalInv={GHR=14b_INV}1454619 Global={GHR=14b}1453273 GlobalInv={GHR=14b}1454619
../fft_3.trace: Total=143398 Local={GHR=12b}134565 LocalInv={GHR=12b_INV}134658 Global={GHR=14b}134416 GlobalInv={GHR=14b}134585
../fft_4.trace: Total=307240 Local={GHR=12b}289820 LocalInv={GHR=12b_INV}290022 Global={GHR=14b}289815 GlobalInv={GHR=14b}289891
../ispell_1.trace: Total=892998 Local={GHR=12b}846648 LocalInv={GHR=12b_INV}846863 Global={GHR=14b}844978 GlobalInv={GHR=14b}845236
../ispell_2.trace: Total=84224869 Local={GHR=14b}81094140 LocalInv={GHR=14b_INV}81094921 Global={GHR=14b}81094140 GlobalInv={GHR=14b}81094921
../jpeg_1.trace: Total=10468313 Local={GHR=14b}9912855 LocalInv={GHR=14b_INV}9913556 Global={GHR=14b}9912855 GlobalInv={GHR=14b}9913556
../jpeg_2.trace: Total=1069396 Local={GHR=14b}1023442 LocalInv={GHR=14b_INV}1024002 Global={GHR=14b}1023442 GlobalInv={GHR=14b}1024002
../jpeg_3.trace: Total=2975508 Local={GHR=14b}2762802 LocalInv={GHR=14b_INV}2763445 Global={GHR=14b}2762802 GlobalInv={GHR=14b}2763445
../jpeg_4.trace: Total=293615 Local={GHR=14b}275372 LocalInv={GHR=14b_INV}275924 Global={GHR=14b}275372 GlobalInv={GHR=14b}275924
../patricia_1.trace: Total=402418 Local={GHR=14b}375233 LocalInv={GHR=14b_INV}375253 Global={GHR=14b}375233 GlobalInv={GHR=14b}375253
../patricia_2.trace: Total=2507039 Local={GHR=14b}2353816 LocalInv={GHR=14b_INV}2353835 Global={GHR=14b}2353816 GlobalInv={GHR=14b}2353835
../qsort_1.trace: Total=150437 Local={GHR=5b}105094 LocalInv={GHR=5b_INV}105096 Global={GHR=14b}105072 GlobalInv={GHR=14b}105074
../qsort_2.trace: Total=1728620 Local={GHR=14b}1370008 LocalInv={GHR=2b_INV}1377615 Global={GHR=14b}1370008 GlobalInv={GHR=14b}1370010
../rijndael_1.trace: Total=117270 Local={GHR=4b}117192 LocalInv={GHR=4b_INV}117215 Global={GHR=14b}117157 GlobalInv={GHR=14b}117180
../rijndael_2.trace: Total=97767 Local={GHR=4b}97705 LocalInv={GHR=4b_INV}97715 Global={GHR=14b}97670 GlobalInv={GHR=14b}97680
../rijndael_3.trace: Total=1218168 Local={GHR=4b}1218090 LocalInv={GHR=4b_INV}1218113 Global={GHR=14b}1218055 GlobalInv={GHR=14b}1218078
../rijndael_4.trace: Total=1015182 Local={GHR=4b}1015120 LocalInv={GHR=4b_INV}1015130 Global={GHR=14b}1015085 GlobalInv={GHR=14b}1015095
../stringsearch_1.trace: Total=15743 Local={GHR=1b}15525 LocalInv={GHR=1b_INV}15526 Global={GHR=14b}15364 GlobalInv={GHR=14b}15384
../stringsearch_2.trace: Total=361099 Local={GHR=3b}357484 LocalInv={GHR=3b_INV}357485 Global={GHR=14b}356855 GlobalInv={GHR=14b}356856
../susan-c_1.trace: Total=39122 Local={GHR=9b}37394 LocalInv={GHR=9b_INV}37415 Global={GHR=14b}37175 GlobalInv={GHR=14b}37195
../susan-c_2.trace: Total=1062352 Local={GHR=10b}991386 LocalInv={GHR=10b_INV}991399 Global={GHR=14b}989056 GlobalInv={GHR=14b}989071
../susan-e_1.trace: Total=74427 Local={GHR=10b}72309 LocalInv={GHR=10b_INV}72345 Global={GHR=14b}72074 GlobalInv={GHR=14b}72113
../susan-e_2.trace: Total=2869864 Local={GHR=14b}2673272 LocalInv={GHR=14b_INV}2673408 Global={GHR=14b}2673272 GlobalInv={GHR=14b}2673408
../susan-s_1.trace: Total=1756847 Local={GHR=14b}1749474 LocalInv={GHR=14b_INV}1749498 Global={GHR=14b}1749474 GlobalInv={GHR=14b}1749498
../susan-s_2.trace: Total=26878317 Local={GHR=14b}26767324 LocalInv={GHR=14b_INV}26767360 Global={GHR=14b}26767324 GlobalInv={GHR=14b}26767360
GLOBAL Total=485419196 Global={GHR=14b}453392019 GlobalInv={GHR=14b}453398851

Appendix C

Totals Tables

| Total | 1 | 2 | 4 | 8 | 16 | 32 | 64 | 128 | 256 | 512 | 1024 | 2048 | 4096 | 8192 | 16384 | 32768 | 65536 | 131072 | 262144 | 524288 | 1048576 |
|----------------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|
| mcf_localINV | 75.3771 | 73.7892 | 73.7836 | 82.8783 | 86.8661 | 88.9726 | 89.5704 | 89.9625 | 90.1964 | 90.4465 | 90.5384 | 90.6394 | 90.6838 | 91.0934 | 91.1807 | 91.2250 | 91.2379 | 91.2604 | 91.2477 | 91.2267 | 91.1696 |
| mcf_gselectINV | 74.9432 | 78.0754 | 82.3712 | 83.0421 | 83.9568 | 86.1686 | 86.7982 | 88.4353 | 91.3616 | 92.0042 | 92.5640 | 92.8002 | 93.2800 | 93.4190 | 93.5195 | 93.5883 | 93.6251 | 93.6692 | 93.6989 | 93.7573 | 93.8097 |
| biasbit | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 |
| agreeINV | 82.0439 | 83.4887 | 85.0784 | 85.5810 | 87.9573 | 89.4145 | 89.6121 | 90.2109 | 91.6998 | 92.4791 | 92.8221 | 92.9810 | 93.5019 | 93.6274 | 93.6649 | 93.6954 | 93.7448 | 93.7859 | 93.8531 | 93.8995 | 93.9540 |
| saturating | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 | 51.1047 |
| mcf_lgshareINV | 76.4629 | 76.1264 | 78.0627 | 78.8600 | 85.1923 | 88.1596 | 90.8477 | 91.4263 | 91.8365 | 92.1489 | 92.3470 | 92.7053 | 93.0237 | 93.1237 | 93.5597 | 93.6477 | 93.7008 | 93.7333 | 93.7735 | 93.8027 | 93.8560 |
| forwards | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 | 44.6228 |
| mcf_bimodal | 61.9786 | 65.6260 | 73.7827 | 77.8175 | 79.1724 | 81.6791 | 82.2127 | 82.4809 | 82.5174 | 82.6058 | 82.6304 | 82.6469 | 82.6665 | 82.6665 | 82.6665 | 82.6665 | 82.6665 | 82.6665 | 82.6665 | 82.6665 | 82.6665 |
| mcf_bimodalINV | 74.9432 | 79.5160 | 83.8119 | 85.7386 | 86.5191 | 86.7151 | 86.9256 | 86.9986 | 87.0323 | 87.0929 | 87.1108 | 87.0956 | 87.1044 | 87.1044 | 87.1044 | 87.1044 | 87.1044 | 87.1044 | 87.1044 | 87.1044 | 87.1044 |
| mcf_global | 64.9945 | 69.6400 | 74.5478 | 77.9639 | 80.4679 | 83.9295 | 85.6532 | 88.3845 | 91.3351 | 91.9919 | 92.5561 | 92.7932 | 93.2730 | 93.4122 | 93.4897 | 93.5274 | 93.5864 | 93.6185 | 93.6928 | 93.7514 | 93.7993 |
| mcf_globalINV | 74.3559 | 75.9690 | 77.2908 | 79.1479 | 81.1101 | 84.0524 | 85.7630 | 88.4353 | 91.3616 | 92.0042 | 92.5640 | 92.8002 | 93.2800 | 93.4190 | 93.4960 | 93.5336 | 93.5924 | 93.6244 | 93.6989 | 93.7573 | 93.8051 |
| saturatingINV | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 | 77.4237 |
| cascadedINV | 76.0160 | 80.0502 | 81.3351 | 82.3712 | 84.1174 | 86.0411 | 86.1026 | 87.2753 | 89.3219 | 91.0219 | 91.9623 | 92.5651 | 92.8277 | 93.4036 | 93.5209 | 93.6141 | 93.6388 | 93.6932 | 93.7321 | 93.8025 | 93.8455 |
| correlated | 62.5056 | 71.9947 | 76.9595 | 79.7833 | 81.8945 | 82.9148 | 83.3611 | 83.8855 | 84.5982 | 86.4069 | 86.5458 | 86.6617 | 86.8122 | 88.4581 | 88.5044 | 88.5820 | 90.3057 | 90.3382 | 90.3621 | 90.3711 | 90.7637 |
| all | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 | 60.7144 |
| mcf_gshareINV | 80.0502 | 81.3351 | 82.3712 | 84.1174 | 86.0411 | 86.1026 | 87.2753 | 89.3219 | 91.0219 | 91.9623 | 92.5651 | 92.8277 | 93.4036 | 93.5209 | 93.6141 | 93.6388 | 93.6932 | 93.7321 | 93.8025 | 93.8455 | 93.8938 |
| mcf_gshare | 66.2711 | 72.0610 | 76.7617 | 80.6244 | 83.8474 | 85.0720 | 87.2227 | 89.2781 | 91.0084 | 91.9547 | 92.5618 | 92.8253 | 93.4022 | 93.5195 | 93.6133 | 93.6379 | 93.6922 | 93.7310 | 93.8014 | 93.8444 | 93.8927 |
| bimodeINV | 74.0397 | 74.5791 | 78.3908 | 81.7353 | 84.3285 | 87.3155 | 89.2738 | 90.3096 | 90.4567 | 90.9678 | 92.3522 | 92.7535 | 93.0296 | 93.1497 | 93.5889 | 93.6472 | 93.6946 | 93.7099 | 93.7470 | 93.7783 | 93.8376 |
| mcf_gselect | 64.9945 | 70.6590 | 74.9782 | 78.4552 | 82.4605 | 84.8856 | 86.5626 | 88.3845 | 91.3351 | 91.9919 | 92.5561 | 92.7932 | 93.2730 | 93.4122 | 93.5160 | 93.5851 | 93.6224 | 93.6678 | 93.6953 | 93.7514 | 93.8068 |
| allINV | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 |
| cascaded | 59.6376 | 66.2711 | 72.0610 | 76.7617 | 80.6244 | 83.8474 | 85.0720 | 87.2227 | 89.2781 | 91.0084 | 91.9547 | 92.5618 | 92.8253 | 93.4022 | 93.5195 | 93.6133 | 93.6379 | 93.6922 | 93.7310 | 93.8014 | 93.8444 |
| backwards | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 | 55.3772 |
| mcf_lgshare | 54.6688 | 63.5622 | 71.3504 | 75.2288 | 84.2848 | 87.6652 | 90.5496 | 91.2221 | 91.8273 | 92.1448 | 92.3452 | 92.7044 | 93.0230 | 93.1233 | 93.5595 | 93.6473 | 93.7006 | 93.7331 | 93.7731 | 93.8023 | 93.8552 |
| cascaded2 | 56.7124 | 64.9945 | 70.6590 | 74.9782 | 78.4552 | 82.4605 | 84.8856 | 86.5626 | 88.3845 | 91.3351 | 91.9919 | 92.5561 | 92.7932 | 93.2730 | 93.4122 | 93.5160 | 93.5851 | 93.6224 | 93.6678 | 93.6953 | 93.7514 |
| mcf_local | 60.8970 | 66.4562 | 70.0944 | 81.9418 | 86.2159 | 88.5929 | 89.2339 | 89.9400 | 90.1876 | 90.4408 | 90.5361 | 90.6370 | 90.6825 | 91.0314 | 91.1776 | 91.2213 | 91.2337 | 91.2554 | 91.2420 | 91.2202 | 91.1621 |
| comb | 85.1895 | 85.7417 | 87.0841 | 88.3190 | 90.5118 | 91.2985 | 91.9365 | 92.3225 | 92.6261 | 92.8230 | 93.0910 | 93.3277 | 93.6768 | 93.8913 | 93.9549 | 94.1177 | 94.2503 | 94.3075 | 94.3489 | 94.4086 | 94.4315 |
| markov | 00.0000 | 00.0000 | 60.2871 | 59.7945 | 68.9840 | 72.5432 | 75.5870 | 77.3652 | 78.9483 | 82.7754 | 83.8767 | 87.1089 | 91.1668 | 91.9461 | 92.6062 | 92.8734 | 93.3663 | 93.5002 | 93.6269 | 93.6796 | 93.7303 |
| backwardsINV | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 |
| forwardsINV | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 |
| markovINV | 00.0000 | 00.0000 | 78.7177 | 77.5658 | 77.5620 | 77.4169 | 78.1832 | 78.7335 | 79.6828 | 82.9703 | 84.0141 | 87.1884 | 91.2133 | 91.9671 | 92.6237 | 92.8843 | 93.3769 | 93.5103 | 93.6367 | 93.6887 | 93.7389 |
| biasbitINV | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 | 83.2705 |
| correlatedINV | 77.6601 | 80.7426 | 82.4548 | 82.9105 | 84.1133 | 84.7563 | 85.0636 | 86.1415 | 87.1670 | 88.8024 | 88.7250 | 89.8201 | 89.9049 | 90.1734 | 90.2074 | 90.3563 | 90.4282 | 90.6526 | 90.6976 | 90.7067 | 90.7941 |
| bimode | 55.8111 | 64.4524 | 71.9158 | 77.3369 | 82.6073 | 87.0564 | 89.0636 | 90.2391 | 90.3951 | 90.9611 | 92.3512 | 92.7523 | 93.0275 | 93.1488 | 93.5884 | 93.6468 | 93.6941 | 93.7097 | 93.7468 | 93.7781 | 93.8372 |
| agree | 80.5150 | 81.7587 | 83.6741 | 85.5393 | 87.9026 | 89.3537 | 89.5517 | 90.2044 | 91.6975 | 92.4767 | 92.8196 | 92.9792 | 93.5013 | 93.6268 | 93.6649 | 93.6953 | 93.7447 | 93.7858 | 93.8531 | 93.8995 | 93.9540 |
| cascaded2INV | 76.0160 | 74.9432 | 78.0754 | 82.3712 | 83.0421 | 83.9568 | 86.1686 | 86.7982 | 88.4353 | 91.3616 | 92.0042 | 92.5640 | 92.8002 | 93.2800 | 93.4190 | 93.5195 | 93.5883 | 93.6251 | 93.6692 | 93.6989 | 93.7573 |

| Average | 1 | 2 | 4 | 8 | 16 | 32 | 64 | 128 | 256 | 512 | 1024 | 2048 | 4096 | 8192 | 16384 | 32768 | 65536 | 131072 | 262144 | 524288 | 1048576 |
|----------------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|
| mcf_localINV | 85.9206 | 85.0225 | 85.6168 | 88.5094 | 90.2406 | 91.0450 | 91.5981 | 91.9445 | 92.1489 | 92.3397 | 92.4227 | 92.5778 | 92.6881 | 93.0563 | 93.2145 | 93.1937 | 93.1361 | 93.0722 | 92.9926 | 92.9038 | 92.8012 |
| mcf_gselectINV | 86.6542 | 88.3040 | 89.3653 | 90.3236 | 90.7957 | 91.3272 | 91.6126 | 92.0276 | 92.5345 | 92.8476 | 93.1119 | 93.2160 | 93.9252 | 93.9702 | 94.0380 | 94.0779 | 94.1024 | 94.1271 | 94.1367 | 94.1644 | 94.1785 |
| biasbit | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 |
| agreeINV | 89.3955 | 89.9111 | 90.5929 | 90.9706 | 91.5623 | 91.9345 | 92.1079 | 92.4163 | 92.7697 | 93.0555 | 93.2435 | 93.3324 | 94.0573 | 94.1158 | 94.1740 | 94.1945 | 94.2272 | 94.2545 | 94.2731 | 94.3066 | 94.3243 |
| saturating | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 | 66.8440 |
| mcf_lgshareINV | 84.9636 | 86.2663 | 86.8626 | 88.0736 | 89.9411 | 91.1946 | 91.9398 | 92.4403 | 92.8754 | 93.0231 | 93.1526 | 93.3460 | 93.5369 | 93.5928 | 94.1514 | 94.2248 | 94.2603 | 94.2814 | 94.3035 | 94.3283 | 94.3497 |
| forwards | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 | 28.0643 |
| mcf_bimodal | 76.4755 | 80.5623 | 85.3620 | 87.2265 | 87.8097 | 89.4616 | 89.7467 | 89.8706 | 89.8892 | 89.9248 | 89.9339 | 89.9382 | 89.9452 | 89.9452 | 89.9452 | 89.9452 | 89.9452 | 89.9452 | 89.9452 | 89.9452 | 89.9452 |
| mcf_bimodalINV | 86.0273 | 87.6182 | 89.3736 | 90.1308 | 90.4250 | 90.5923 | 90.6762 | 90.7271 | 90.7458 | 90.7698 | 90.7764 | 90.7750 | 90.7781 | 90.7781 | 90.7781 | 90.7781 | 90.7781 | 90.7781 | 90.7781 | 90.7781 | 90.7781 |
| mcf_global | 79.2236 | 82.5123 | 85.6461 | 87.2722 | 88.3970 | 89.5989 | 90.4147 | 91.3811 | 92.2254 | 92.5560 | 92.8408 | 92.9382 | 93.6476 | 93.7164 | 93.7516 | 93.7249 | 93.7255 | 93.7002 | 93.6888 | 93.5965 | 93.5692 |
| mcf_globalINV | 85.2786 | 86.0329 | 86.8439 | 88.0462 | 88.9021 | 89.7993 | 90.5559 | 91.4691 | 92.2903 | 92.6110 | 92.8968 | 92.9943 | 93.7099 | 93.7791 | 93.8089 | 93.7802 | 93.7811 | 93.7373 | 93.7207 | 93.6293 | 93.6073 |
| saturatingINV | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 | 84.7671 |
| cascadedINV | 86.6652 | 87.5584 | 88.4042 | 89.5727 | 89.9827 | 90.8082 | 91.1607 | 91.5559 | 91.9015 | 92.4242 | 92.7622 | 93.0045 | 93.1302 | 93.8736 | 93.9234 | 93.9892 | 94.0106 | 94.0374 | 94.0592 | 94.0815 | 94.1066 |
| correlated | 79.7513 | 86.5164 | 88.3081 | 88.9700 | 89.8663 | 90.3169 | 90.6938 | 90.8313 | 91.2108 | 91.7800 | 91.8916 | 91.9471 | 91.9941 | 92.4152 | 92.4606 | 92.5374 | 92.5527 | 92.7318 | 92.7588 | 92.7735 | 92.8159 |
| all | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 | 63.6663 |
| mcf_gshareINV | 87.5584 | 88.4042 | 89.5727 | 89.9827 | 90.8082 | 91.1607 | 91.5559 | 91.9015 | 92.4242 | 92.7622 | 93.0045 | 93.1302 | 93.8736 | 93.9234 | 93.9892 | 94.0106 | 94.0374 | 94.0592 | 94.0815 | 94.1066 | 94.1218 |
| mcf_gshare | 82.4856 | 85.7894 | 87.7597 | 88.9896 | 90.1719 | 90.7510 | 91.3855 | 91.8344 | 92.3776 | 92.7238 | 92.9671 | 93.0941 | 93.8364 | 93.8869 | 93.9579 | 93.9787 | 94.0059 | 94.0277 | 94.0500 | 94.0751 | 94.0904 |
| bimodeINV | 83.6578 | 85.5811 | 87.6038 | 89.2047 | 90.1921 | 90.9978 | 91.7108 | 92.0996 | 92.3311 | 92.5761 | 92.9097 | 93.1042 | 93.2655 | 93.3350 | 94.0439 | 94.0896 | 94.1487 | 94.1632 | 94.1895 | 94.2146 | 94.2281 |
| mcf_gselect | 80.9430 | 84.6235 | 87.4264 | 89.2887 | 90.0849 | 90.9081 | 91.3480 | 91.8957 | 92.4720 | 92.7952 | 93.0567 | 93.1602 | 93.8763 | 93.9294 | 93.9977 | 94.0447 | 94.0699 | 94.0945 | 94.1047 | 94.1328 | 94.1468 |
| allINV | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 |
| cascaded | 74.4838 | 82.4856 | 85.7894 | 87.7597 | 88.9896 | 90.1719 | 90.7510 | 91.3855 | 91.8344 | 92.3776 | 92.7238 | 92.9671 | 93.0941 | 93.8364 | 93.8869 | 93.9579 | 93.9787 | 94.0059 | 94.0277 | 94.0500 | 94.0751 |
| backwards | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 | 71.9357 |
| mcf_lgshare | 65.8841 | 77.3192 | 84.4998 | 86.8143 | 89.4140 | 90.9426 | 91.8439 | 92.3918 | 92.8550 | 93.0043 | 93.1363 | 93.3302 | 93.5166 | 93.5763 | 94.1363 | 94.2083 | 94.2452 | 94.2680 | 94.2894 | 94.3131 | 94.3347 |
| cascaded2 | 71.7842 | 80.9430 | 84.6235 | 87.4264 | 89.2887 | 90.0849 | 90.9081 | 91.3480 | 91.8957 | 92.4720 | 92.7952 | 93.0567 | 93.1602 | 93.8763 | 93.9294 | 93.9977 | 94.0447 | 94.0699 | 94.0945 | 94.1047 | 94.1328 |
| mcf_local | 75.1933 | 82.2155 | 84.4913 | 88.0172 | 89.9202 | 90.8649 | 91.4362 | 91.9107 | 92.1283 | 92.3225 | 92.3134 | 92.4740 | 92.6549 | 92.8340 | 93.1531 | 93.1211 | 93.0558 | 92.9799 | 92.8901 | 92.7915 | 92.6768 |
| comb | 90.9173 | 91.3116 | 91.7429 | 92.1669 | 92.4395 | 92.7505 | 93.0845 | 93.2811 | 93.4142 | 93.5849 | 93.7498 | 93.8298 | 94.3887 | 94.4334 | 94.5259 | 94.5588 | 94.5835 | 94.6070 | 94.6185 | 94.6540 | 94.6694 |
| markov | 0.0000 | 0.0000 | 70.1381 | 69.8547 | 78.5168 | 82.6024 | 86.3641 | 87.5333 | 88.4622 | 89.5850 | 90.2944 | 91.2792 | 92.2460 | 92.6084 | 92.8955 | 93.0058 | 93.7242 | 93.8116 | 93.8614 | 93.9082 | 93.9642 |
| backwardsINV | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 |
| forwardsINV | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 |
| markovINV | 0.0000 | 0.0000 | 85.3572 | 85.9616 | 86.4885 | 87.3833 | 87.9013 | 88.4251 | 89.0776 | 89.8125 | 90.4408 | 91.3906 | 92.3276 | 92.6743 | 92.9619 | 93.0676 | 93.7862 | 93.8737 | 93.9225 | 93.9681 | 94.0150 |
| biasbitINV | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 | 89.4867 |
| correlatedINV | 87.4518 | 89.1370 | 89.6576 | 89.9887 | 90.5171 | 90.7965 | 90.9940 | 91.2957 | 91.6845 | 91.9317 | 92.0087 | 92.0589 | 92.1057 | 92.5260 | 92.5694 | 92.7741 | 92.7984 | 92.8469 | 92.8650 | 92.8766 | 92.9240 |
| bimode | 66.8923 | 79.0117 | 84.4848 | 87.8724 | 89.5504 | 90.6512 | 91.4151 | 91.9856 | 92.2384 | 92.5256 | 92.8879 | 93.0854 | 93.2440 | 93.3153 | 94.0238 | 94.0697 | 94.1291 | 94.1448 | 94.1729 | 94.1979 | 94.2114 |
| agree | 89.0600 | 89.5641 | 90.2795 | 90.7960 | 91.3893 | 91.7944 | 92.0362 | 92.3862 | 92.7513 | 93.0370 | 93.2209 | 93.3110 | 94.0354 | 94.0945 | 94.1587 | 94.1794 | 94.2132 | 94.2404 | 94.2591 | 94.2926 | 94.3103 |
| cascaded2INV | 86.2683 | 86.6542 | 88.3040 | 89.3653 | 90.3236 | 90.7957 | 91.3272 | 91.6126 | 92.0276 | 92.5345 | 92.8476 | 93.1119 | 93.2160 | 93.9252 | 93.9702 | 94.0380 | 94.0779 | 94.1024 | 94.1271 | 94.1367 | 94.1644 |